

22p_9252_tazmeen_afroz_canny

October 6, 2025

```
[ ]: import cv2
```

```
import numpy as np
```

```
from scipy import ndimage
```

```
import matplotlib.pyplot as plt
```

```
[ ]: def convolution2d(image, kernel):
```

```
    image_h, image_w = image.shape
```

```
    kernel_h, kernel_w = kernel.shape
```

```
    padding_h = (kernel_h - 1) // 2
```

```
    padding_w = (kernel_w - 1) // 2
```

```
    padded_image = np.pad(image, ((padding_h, padding_h), (padding_w, padding_w)), mode='constant', constant_values=0)
```

```
    output = np.zeros((image_h, image_w))
```

```
    print(output.size)
```

```
    for i in range(image_h):
```

```
        for j in range(image_w):
```

```
            new = padded_image[i:i+kernel_h, j:j+kernel_w]
```

```
            output[i, j] = np.sum(new * kernel)
```

```
    return output
```

```
image = np.random.rand(8, 4)
```

```
print(image)
```

```
kernel = np.random.rand(3, 3)
```

```
kernel = np.fliplr(np.fliplr(kernel))
```

```
convolution2d(image, kernel)
```

```
[[0.15732629 0.61268274 0.14615 0.22660716]
```

```
[0.25869883 0.46955753 0.1171389 0.44406555]
```

```
[0.46716219 0.16372021 0.93813416 0.91465293]
```

```
[0.07879303 0.66933523 0.80987996 0.08128713]
```

```
[0.71368074 0.75670111 0.17191526 0.25466064]
```

```
[0.58678824 0.72518546 0.62166664 0.33822399]
```

```
[0.81172379 0.09778313 0.17076208 0.5300113 ]  
[0.22321054 0.53610429 0.41620201 0.73945046]]
```

32

```
[ ]: array([[0.94623699, 1.11559242, 1.44219021, 0.54423524],  
           [1.14554323, 1.82766166, 2.11603252, 1.49228617],  
           [1.12470879, 2.323      , 2.72426519, 2.30691035],  
           [1.47772073, 2.63790457, 2.73933608, 1.62200133],  
           [2.03085546, 2.92901178, 2.42419639, 1.12360946],  
           [1.94980052, 2.78286481, 2.19276048, 1.28587536],  
           [1.5354613 , 2.21707997, 2.07604941, 1.37987655],  
           [0.93565385, 1.28173656, 1.74405339, 1.17617656]])
```

```
[ ]: def canny_edge_detection(image, low=50, high=150):  
    # Step 1: blurring  
  
    # step 2: edges  
    # sobel used for both blurring and edge detection  
    sobel_x = np.array([[-1, 0, 1], [-2, 0, 2], [-1, 0, 1]])  
    sobel_y = np.array([[-1, -2, -1], [0, 0, 0], [1, 2, 1]])  
  
    canny_gx = convolution2d(image.astype(np.float32), sobel_x)  
    canny_gy = convolution2d(image.astype(np.float32), sobel_y)  
    #finding magnitude  
    combined_canny = np.sqrt(canny_gx**2 + canny_gy**2)  
  
    # arctan2 always gives angle in range of -180 to 180 degrees so no need to  
    ↪check for angles higher than 180 here  
    angle = np.arctan2(canny_gy, canny_gx)  
    angle[angle < 0] += np.pi  
  
    # Step 3: Non-maximum suppression  
    nms_r = nms(combined_canny, angle)  
  
    # Step 4: Hysteresis thresholding  
    edges = hysteresis(nms_r, low, high)  
  
    return edges  
  
def nms(combined_canny, angle):  
    height, width = combined_canny.shape  
    suppressed = np.zeros((height, width), dtype=np.float64)  
  
    for i in range(1, height - 1):  
        for j in range(1, width - 1):
```

```

current_magnitude = combined_canny[i, j]
current_angle = angle[i, j]

# Use edge direction (perpendicular to gradient): add pi/2
edge_angle = current_angle + np.pi / 2
if edge_angle > np.pi:
    edge_angle -= np.pi
#dx and dy are the offsets to check the neighboring pixels in the
#edge direction
#dx gives the unit step in x direction and dy gives the unit step
#in y direction
dx = int(np.cos(edge_angle))
dy = int(np.sin(edge_angle))

# Coordinates of the two neighboring pixels in the edge direction
# one step ahead and one step behind along the axis of the edge

x1 = j + dx    # one step ahead in x direction from current pixel
y1 = i + dy    # one step ahead in y direction from current pixel
x2 = j - dx    # one step behind in x direction from current pixel
y2 = i - dy    # one step behind in y direction from current pixel

# Check if the neighboring pixels are within image bounds
if (0 <= x1 < width and 0 <= y1 < height and
    0 <= x2 < width and 0 <= y2 < height):

    # calculate the magnitudes of the neighboring pixels
    neighbor1_mag = combined_canny[y1, x1]
    neighbor2_mag = combined_canny[y2, x2]

    # Suppress if current pixel is not greater than both neighbors
    if (current_magnitude >= neighbor1_mag and
        current_magnitude >= neighbor2_mag):
        suppressed[i, j] = current_magnitude
return suppressed

def hysteresis(nms_r, low, high):
    """Hysteresis thresholding with weak edges promoted to high threshold
    value"""
    height, width = nms_r.shape

    strong_edges = (nms_r >= high)
    weak_edges = ((nms_r >= low) & (nms_r < high))

    edges = np.zeros((height, width), dtype=np.uint8)
    edges[strong_edges] = high

```

```

changed = True
iterations = 0
while changed and iterations < 10:
    changed = False
    iterations += 1

    for i in range(1, height - 1):
        for j in range(1, width - 1):
            if weak_edges[i, j] and edges[i, j] == 0:
                neighborhood = edges[i-1:i+2, j-1:j+2]
                if np.any(neighborhood >= high): # check connected to ↵
                    "strong"
                    edges[i, j] = high
                    changed = True
return edges

img = cv2.imread('/content/phool.png', cv2.IMREAD_GRAYSCALE)

# Custom Canny
edges_custom = canny_edge_detection(img, low=50, high=255)


# OpenCV Canny
edges_cv = cv2.Canny(img, 50, 100)

plt.figure(figsize=(14,6))
plt.subplot(1,2,1)
plt.imshow(edges_custom, cmap='gray')
plt.title('Custom Canny')
plt.axis('off')
plt.subplot(1,2,2)
plt.imshow(edges_cv, cmap='gray')
plt.title('OpenCV Canny')
plt.axis('off')
plt.show()

```

155694

155694



code with interpolation

The integer version implemented above rounds the edge direction to the nearest horizontal, vertical, or diagonal pixel, which makes the algorithm simpler but less accurate, since many true edge directions lie between these fixed angles. This often produces thicker or denser edges because suppression is not perfectly aligned with the gradient. The interpolation version instead follows the exact gradient direction by sampling magnitudes at fractional coordinates using bilinear interpolation, which estimates values between real pixels. This makes edge thinning more precise and produces smoother, cleaner edges that closely match the results of OpenCV's Canny detector.

```
[ ]: # code with interpolation

def bilinear_interpolate(img, y, x):

    h, w = img.shape
    if x < 0 or x >= w - 1 or y < 0 or y >= h - 1:
        return 0

    x1, y1 = int(np.floor(x)), int(np.floor(y))
    x2, y2 = x1 + 1, y1 + 1
    fx, fy = x - x1, y - y1

    return (img[y1, x1] * (1 - fx) * (1 - fy) +
            img[y1, x2] * fx * (1 - fy) +
            img[y2, x1] * (1 - fx) * fy +
            img[y2, x2] * fx * fy)

def non_max_suppression(magnitude, angle):
    """Thin edges using gradient direction + interpolation"""
    h, w = magnitude.shape
    suppressed = np.zeros((h, w), dtype=np.float32)
```

```

for i in range(1, h - 1):
    for j in range(1, w - 1):
        mag = magnitude[i, j]
        if mag == 0:
            continue

        edge_angle = angle[i, j] + np.pi / 2
        dx, dy = np.cos(edge_angle), np.sin(edge_angle)

        # Interpolated magnitudes in both directions
        mag1 = bilinear_interpolate(magnitude, i + dy, j + dx)
        mag2 = bilinear_interpolate(magnitude, i - dy, j - dx)

        if mag >= mag1 and mag >= mag2:
            suppressed[i, j] = mag

return suppressed

def hysteresis(nms, low, high):
    """Hysteresis thresholding"""
    h, w = nms.shape
    strong = (nms >= high)
    weak = ((nms >= low) & (nms < high))

    edges = np.zeros((h, w), dtype=np.uint8)
    edges[strong] = 255

    changed = True
    while changed:
        changed = False
        for i in range(1, h - 1):
            for j in range(1, w - 1):
                if weak[i, j] and edges[i, j] == 0:
                    if np.any(edges[i-1:i+2, j-1:j+2] == 255):
                        edges[i, j] = 255
                        changed = True

    return edges

def custom_canny(image, low=50, high=150):

    # Step 2: Gradient calculation
    sobel_x = np.array([[-1, 0, 1], [-2, 0, 2], [-1, 0, 1]], dtype=np.float32)
    sobel_y = np.array([[-1, -2, -1], [0, 0, 0], [1, 2, 1]], dtype=np.float32)

```

```

gx = convolution2d(image.astype(np.float32), sobel_x)
gy = convolution2d(image.astype(np.float32), sobel_y)

magnitude = np.sqrt(gx**2 + gy**2)
angle = np.arctan2(gy, gx)
angle[angle < 0] += np.pi

# Step 3: Non-maximum suppression
nms = non_max_suppression(magnitude, angle)

# Step 4: Hysteresis
edges = hysteresis(nms, low, high)
return edges

img = cv2.imread("/content/phool.png", cv2.IMREAD_GRAYSCALE)

edges_custom = custom_canny(img, low=100, high=150)
edges_cv = cv2.Canny(img, 100, 150)

plt.figure(figsize=(14, 6))



plt.subplot(1, 2, 1)
plt.imshow(edges_custom, cmap="gray")
plt.title("Custom Canny"); plt.axis("off")

plt.subplot(1, 2, 2)
plt.imshow(edges_cv, cmap="gray")
plt.title("OpenCV Canny"); plt.axis("off")

plt.show()

```

155694
155694



[]: