



# Computer Vision

---

Dr. Muhammad Tahir

DEPARTMENT OF COMPUTER SCIENCE,  
FAST-NUCES, Peshawar

# Course Details

---

**LECTURES:** Monday  
& Wednesday

---

**TIMINGS:**  
9:30 am – 11:00 am

---



**MY OFFICE:**

**OFFICE HOURS:**

---

**EMAIL:** [m.tahir@nu.edu.pk](mailto:m.tahir@nu.edu.pk)

---



# References

The material in these slides are based on:

1

Rick Szeliski's book: [Computer Vision: Algorithms and Applications](#)

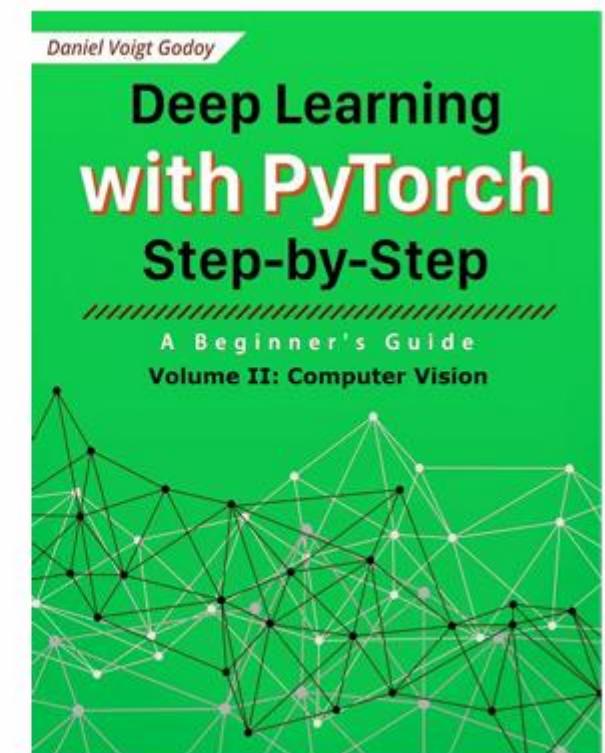
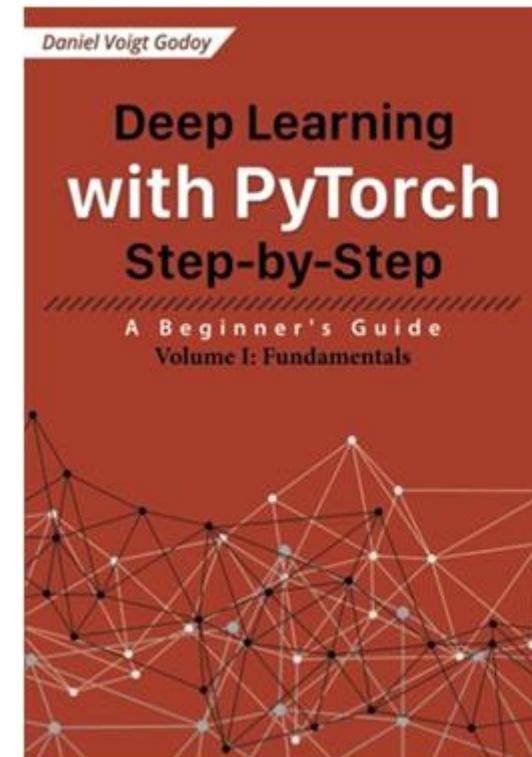
2

Forsythe and Ponce: [Computer Vision: A Modern Approach](#)

# Recommended Books

---

Deep Learning with PyTorch Step-by-Step by Daniel Voigt Godoy



# Course Learning Outcomes

---

No	CLO	Domain	Taxonomy Level	PLO
1	Understand the view geometry concepts, multi-scale representation, edge detection and detection of other primitives, stereo, motion and object recognition.	Cognitive		
2	Assess which methods to use for solving a given problem, and analyse the accuracy of the methods Skills	Cognitive		
3	Apply appropriate image processing methods for image filtering, image restoration, image reconstruction, segmentation, classification and representation	Cognitive		



# Outline

---

Deep Learning Fundamentals

---

[Deep Learning Tutorial - GeeksforGeeks](#)

---

---

---

---

# AlexNet

# AlexNet

---

- Classify a  $224 \times 224$  RGB image into 1000 ImageNet classes.
- Five convolutional stages → three fully-connected (FC) layers → softmax.
- ReLU (instead of tanh), **overlapping** max-pooling, **data augmentation** + heavy dropout, and training split across **two GPUs** with “grouped” convolutions.
- ~61M parameters (most are in the two 4096-unit FC layers).

# AlexNet Architecture ( $224 \times 224$ Input Version)

---

## Stage 1

1. **Conv1:** 96 filters, **11×11**, stride **4**, padding **2**


Output:  $55 \times 55 \times 96$

Activation: **ReLU**

Followed by **LRN** (local response normalization)

2. **MaxPool1:**  $3 \times 3$ , stride 2 (overlapping pooling)

Output:  $27 \times 27 \times 96$



Note: Standard AlexNet used  $227 \times 227$  input size

# AlexNet Architecture ( $224 \times 224$ Input Version)

---

## Stage 2


3. **Conv2:** 256 filters, **5×5**, stride 1, **pad 2**

Output:  $27 \times 27 \times 256$

Activation: ReLU

**LRN**

4. **MaxPool2:**  $3 \times 3$ , stride 2 →  $13 \times 13 \times 256$



# AlexNet Architecture ( $224 \times 224$ Input Version)


---

## Stage 3

5. **Conv3:** 384 filters, **3×3**, stride 1, pad 1

Output:  $13 \times 13 \times 384$

Activation: ReLU



# AlexNet Architecture ( $224 \times 224$ Input Version)


---

## Stage 4

6. **Conv4:** 384 filters, **3×3**, stride 1, pad 1

Output:  $13 \times 13 \times 384$

Activation: ReLU



# AlexNet Architecture ( $224 \times 224$ Input Version)


---

## Stage 5

7. **Conv5:** 256 filters, **3×3**, stride 1, pad 1

Output:  $13 \times 13 \times 256$

Activation: ReLU



8. **MaxPool3:**  $3 \times 3$ , stride 2 →  $6 \times 6 \times 256$

# AlexNet Architecture ( $224 \times 224$ Input Version)

---

## **Flatten Layer**

**Input:**  $6 \times 6 \times 256 = 9216$

→ flattened into a 1D vector of length **9216**

# AlexNet Architecture ( $224 \times 224$ Input Version)

---


## 9. Fully Connected Layer 1

**Input units:** 9216

**Output units:** 4096

**Activation:** ReLU

**Dropout:** 0.5



# AlexNet Architecture ( $224 \times 224$ Input Version)

---


## 10. Fully Connected Layer 2

**Input units:** 4096

**Output units:** 4096

**Activation:** ReLU

**Dropout:** 0.5



# AlexNet Architecture ( $224 \times 224$ Input Version)


---

## 11. Fully Connected Layer 3 (Output Layer)



**Input units:** 4096

**Output units:** 1000

**Activation:** Softmax (for 1000 ImageNet classes)



# AlexNet



# Normalization & activations

---

- **ReLU** dramatically sped up convergence compared to tanh/sigmoid used pre-2012.
- **LRN** (across channels) encouraged competition between features; it's **obsolete** now—**BatchNorm** is better.

# Training setup (2012, from the paper)

---

- **SGD with momentum** 0.9, weight decay 5e-4
- Initial LR ~0.01, dropped  $\times 10$  on plateau.
- **Batch size 128.**
- **Data augmentation:** random crops and horizontal flips; “**PCA color jitter**” (additive RGB lighting noise).
- **Dropout 0.5** on FC6 and FC7 to curb overfitting.

<b>Layer</b>	<b>Type</b>	<b>Filter / Kernel</b>	<b>Stride</b>	<b>Padding</b>	<b>Output Volume</b>	<b>Activation</b>
1	Conv	11×11, 96 filters	4	2	55×55×96	ReLU
2	Max Pool	3×3	2	0	27×27×96	-
3	Conv	5×5, 256 filters	1	2	27×27×256	ReLU
4	Max Pool	3×3	2	0	13×13×256	-
5	Conv	3×3, 384 filters	1	1	13×13×384	ReLU
6	Conv	3×3, 384 filters	1	1	13×13×384	ReLU
7	Conv	3×3, 256 filters	1	1	13×13×256	ReLU
8	Max Pool	3×3	2	0	6×6×256	-
9	FC	-	-	-	4096	ReLU + Dropout
10	FC	-	-	-	4096	ReLU + Dropout
11	FC	-	-	-	1000	Softmax

# AlexNet Architecture ( $224 \times 224$ Input Version)

---

- **Conv1** learns large edge and color blobs (coarse features).
- **Conv2–Conv5** learn increasingly complex features — corners, textures, object parts.
- **FC6 & FC7** act as high-level feature combiners.
- **FC8 (Softmax)** outputs the final classification probabilities.

---

VGG

# VGG-16 Architecture ( $224 \times 224$ Input Version)

---

## Input

Image size:  $224 \times 224 \times 3$

- VGG was designed to show that using **very small filters ( $3 \times 3$ )** repeatedly can be as effective as larger filters (like  $7 \times 7$  or  $11 \times 11$ ), but with fewer parameters and more non-linearities.

# VGG-16 Architecture ( $224 \times 224$ Input Version)

---

## **Conv Block 1**

**Conv1\_1:** 64 filters,  $3 \times 3$ , stride 1, pad 1

**Input:**  $224 \times 224 \times 3$  → **Output:**  $224 \times 224 \times 64$

**Activation:** ReLU

**Conv1\_2:** 64 filters,  $3 \times 3$ , stride 1, pad 1

**Output:**  $224 \times 224 \times 64$

**Activation:** ReLU

**MaxPool1:**  $2 \times 2$ , stride 2

**Output:**  $112 \times 112 \times 64$

# VGG-16 Architecture ( $224 \times 224$ Input Version)

---

## Conv Block 2

**Conv2\_1:** 128 filters,  $3 \times 3$ , stride 1, pad 1

**Input:**  $112 \times 112 \times 64$  → **Output:**  $112 \times 112 \times 128$

**Activation:** ReLU

**Conv2\_2:** 128 filters,  $3 \times 3$ , stride 1, pad 1

**Output:**  $112 \times 112 \times 128$

**Activation:** ReLU

**MaxPool2:**  $2 \times 2$ , stride 2

**Output:**  $56 \times 56 \times 128$

# VGG-16 Architecture ( $224 \times 224$ Input Version)

---

## Conv Block 3

**Conv3\_1:** 256 filters,  $3 \times 3$ , stride 1, pad 1

**Input:**  $56 \times 56 \times 128$  → **Output:**  $56 \times 56 \times 256$

**Activation:** ReLU

**Conv3\_2:** 256 filters,  $3 \times 3$ , stride 1, pad 1

**Output:**  $56 \times 56 \times 256$

**Activation:** ReLU

**Conv3\_3:** 256 filters,  $3 \times 3$ , stride 1, pad 1

**Output:**  $56 \times 56 \times 256$

**Activation:** ReLU

**MaxPool3:**  $2 \times 2$ , stride 2

**Output:**  $28 \times 28 \times 256$

# VGG-16 Architecture ( $224 \times 224$ Input Version)

---

## Conv Block 4

**Conv4\_1:** 512 filters,  $3 \times 3$ , stride 1, pad 1

**Input:**  $28 \times 28 \times 256$  → **Output:**  $28 \times 28 \times 512$

**Activation:** ReLU

**Conv4\_2:** 512 filters,  $3 \times 3$ , stride 1, pad 1

**Output:**  $28 \times 28 \times 512$

**Activation:** ReLU

**Conv4\_3:** 512 filters,  $3 \times 3$ , stride 1, pad 1

**Output:**  $28 \times 28 \times 512$

**Activation:** ReLU

**MaxPool4:**  $2 \times 2$ , stride 2

**Output:**  $14 \times 14 \times 512$

# VGG-16 Architecture ( $224 \times 224$ Input Version)

---

## Conv Block 5

**Conv5\_1:** 512 filters,  $3 \times 3$ , stride 1, pad 1

**Input:**  $14 \times 14 \times 512$  → **Output:**  $14 \times 14 \times 512$

**Activation:** ReLU

**Conv5\_2:** 512 filters,  $3 \times 3$ , stride 1, pad 1

**Output:**  $14 \times 14 \times 512$

**Activation:** ReLU

**Conv5\_3:** 512 filters,  $3 \times 3$ , stride 1, pad 1

**Output:**  $14 \times 14 \times 512$

**Activation:** ReLU

**MaxPool5:**  $2 \times 2$ , stride 2

**Output:**  $7 \times 7 \times 512$

# VGG-16 Architecture ( $224 \times 224$ Input Version)

---

## Fully Connected Layers

**Flatten:**  $7 \times 7 \times 512 = 25088$  features

**FC6:**  $25088 \rightarrow 4096$

**Activation:** ReLU, Dropout (0.5)

**FC7:**  $4096 \rightarrow 4096$

**Activation:** ReLU, Dropout (0.5)

**FC8:**  $4096 \rightarrow 1000$

**Activation:** Softmax (ImageNet classes)

# VGG-16 Architecture ( $224 \times 224$ Input Version)

---

## Training Summary

**Total parameters:**  $\approx 138$  million

**Optimizer:** SGD + Momentum (0.9)

**Batch size:** 256

**Weight decay:**  $5 \times 10^{-4}$

**Training time:**  $\approx 2\text{--}3$  weeks on 4 GPUs

**Data augmentation:** random cropping, horizontal flips, RGB jitter

# VGG-16 Architecture ( $224 \times 224$ Input Version)

---

## Design Philosophy

Only **3×3 filters** (instead of AlexNet's mixed sizes).

**Deeper network** (16 layers with weights).


Each max-pool halves spatial size.

**ReLU after every convolution.**

No LRN — replaced with deeper stacking.

Large parameter count dominated by FC layers.

# VGG-16 Architecture ( $224 \times 224$ Input Version)



- 138 million parameters
- Training: 2-3 week (4 GPUs)

<b>Block</b>	<b>Layers</b>	<b>Output Volume</b>
<b>Input</b>	RGB image	224×224×3
<b>Conv Block 1</b>	2 × (3×3 conv, 64 filters) + MaxPool	112×112×64
<b>Conv Block 2</b>	2 × (3×3 conv, 128 filters) + MaxPool	56×56×128
<b>Conv Block 3</b>	3 × (3×3 conv, 256 filters) + MaxPool	28×28×256
<b>Conv Block 4</b>	3 × (3×3 conv, 512 filters) + MaxPool	14×14×512
<b>Conv Block 5</b>	3 × (3×3 conv, 512 filters) + MaxPool	7×7×512
<b>Flatten</b>	—	25088 (7×7×512)
<b>FC6</b>	Fully Connected (4096) + ReLU + Dropout	4096
<b>FC7</b>	Fully Connected (4096) + ReLU + Dropout	4096
<b>FC8</b>	Fully Connected (1000) + Softmax	1000

---

# ResNet

# ResNet (Residual Network – ResNet-152)

---

**ResNet (Residual Network)** was introduced by **Kaiming He et al., 2015** in the paper "*Deep Residual Learning for Image Recognition*"(arXiv:1512.03385).

- It **revolutionized deep learning** by enabling networks with **over 100 layers** to train effectively — solving the **vanishing gradient problem** using *skip (shortcut) connections*.

# ResNet (Residual Network – ResNet-152)

---

Property	Description
<b>Input image</b>	224×224×3
<b>Layers (ResNet-152)</b>	152 (with 50, 101, and 152 variants)
<b>Parameters</b>	~60 million
<b>Training time</b>	2–3 weeks (on 8 GPUs, 2015 standard)
<b>Building blocks</b>	7×7 conv, 3×3 conv, BatchNorm, Max/Average Pooling, Skip connections
<b>Optimizer</b>	SGD + Momentum

# ResNet (Residual Network – ResNet-152)

---

ResNet comes in two types:

- **Basic Block** – used in shallower models (ResNet-18, ResNet-34)
  - Conv3x3 → BN → ReLU → Conv3x3 → BN → Add skip connection → ReLU
- **Bottleneck Block** – used in deeper models (ResNet-50, ResNet-101, ResNet-152)
  - Conv1x1 → BN → ReLU → Conv3x3 → BN → ReLU → Conv1x1 → BN → Add skip connection → ReLU

The **bottleneck** reduces and then restores the number of channels to make training faster.

# ResNet (Residual Network – ResNet-152)

---

## **Stage 1: Initial Layers**

**Conv1:** 64 filters,  $7 \times 7$ , stride 2, pad 3

**Output:**  $112 \times 112 \times 64$

**Activation:** ReLU

**Batch Normalization** applied

**MaxPool:**  $3 \times 3$ , stride 2

**Output:**  $56 \times 56 \times 64$

# ResNet (Residual Network – ResNet-152)

---

## **Stage 2: Conv2\_x (3 Bottleneck Blocks)**

Each block:

$1 \times 1$ , 64 filters  $\rightarrow$   $3 \times 3$ , 64 filters  $\rightarrow$   $1 \times 1$ , 256 filters

Shortcut: identity or  $1 \times 1$  projection

**Output:**  $56 \times 56 \times 256$

# ResNet (Residual Network – ResNet-152)

---

## **Stage 3: Conv3\_x (8 Bottleneck Blocks)**

Each block:

$1 \times 1$ , 128 filters  $\rightarrow$   $3 \times 3$ , 128 filters  $\rightarrow$   $1 \times 1$ , 512 filters

**First block stride = 2** (reduces spatial size)

**Output:**  $28 \times 28 \times 512$

# ResNet (Residual Network – ResNet-152)

---

## **Stage 4: Conv4\_x (36 Bottleneck Blocks)**

Each block:

$1 \times 1$ , 256 filters  $\rightarrow$   $3 \times 3$ , 256 filters  $\rightarrow$   $1 \times 1$ , 1024 filters

**First block stride = 2**

**Output:**  $14 \times 14 \times 1024$

# ResNet (Residual Network – ResNet-152)

---

## **Stage 5: Conv5\_x (3 Bottleneck Blocks)**

Each block:

$1 \times 1$ , 512 filters  $\rightarrow$   $3 \times 3$ , 512 filters  $\rightarrow$   $1 \times 1$ , 2048 filters

**First block stride = 2**

**Output:**  $7 \times 7 \times 2048$

# ResNet (Residual Network – ResNet-152)

---

## **Global Average Pooling**

Averages each feature map ( $7 \times 7$ ) into a single value.

**Output:**  $1 \times 1 \times 2048$

# ResNet (Residual Network – ResNet-152)

---

## **Fully Connected Layer**

**Input:** 2048

**Output:** 1000

**Activation:** Softmax (ImageNet classes)

# ResNet (Residual Network – ResNet-152)

---

## Residual Block Example (Inside the Network)

Let's illustrate one **bottleneck residual block**:

Layer	Type	Filters	Kernel	Stride	Output Size	Skip Connection
1	Conv + BN + ReLU	64	1×1	1	56×56×64	Identity
2	Conv + BN + ReLU	64	3×3	1	56×56×64	—
3	Conv + BN	256	1×1	1	56×56×256	—
Add	—	—	—	—	56×56×256	( $y = F(x) + x$ )
Output	ReLU	—	—	—	56×56×256	—

# ResNet (Residual Network – ResNet-152)

---

<b>Stage</b>	<b>Output Size</b>	<b>Block Type</b>	<b>#Blocks</b>	<b>Parameters</b>
Conv1	$112 \times 112 \times 64$	$7 \times 7, 64$	1	—
Conv2_x	$56 \times 56 \times 256$	Bottleneck	3	—
Conv3_x	$28 \times 28 \times 512$	Bottleneck	8	—
Conv4_x	$14 \times 14 \times 1024$	Bottleneck	36	—
Conv5_x	$7 \times 7 \times 2048$	Bottleneck	3	—
AvgPool + FC	$1 \times 1 \times 1000$	—	—	—

# ResNet (Residual Network – ResNet-152)

---

Feature	Description
<b>Skip connections</b>	Enable training of >100 layers easily
<b>Batch Normalization</b>	Stabilizes and speeds up convergence
<b>Global Average Pooling</b>	Reduces overfitting (no large FC layers like VGG)
<b>Bottleneck structure</b>	Efficient depth expansion
<b>ReLU everywhere</b>	Non-linearity after each conv
<b>Identity mapping</b>	Keeps features stable through depth

# Why ResNet Succeeded

---

Deep networks can be trained **without degradation**.


Each block learns **only the residual difference**.

Skip connections act like **shortcuts for gradients**.

Model achieved **1st place in ILSVRC 2015** with **3.6% Top-5 error**, a huge leap in performance.

# ResNet

---



Kaiming He, <https://arxiv.org/pdf/1512.03385.pdf>

152 layers!!!

7x7 convolutional layers, 3x3 convolutional layers, batch normalization, max and average pooling.

Parameters: 60 million

Training time: 2-3 weeks (8 GPUs)

# 1x1 Convolution

---

**Input feature map:**  $3 \times 3 \times 3$ (Height  $\times$  Width  $\times$  Channels)

$$R = \begin{bmatrix} 1 & 2 & 3 \\ 4 & 5 & 6 \\ 7 & 8 & 9 \end{bmatrix}, G = \begin{bmatrix} 2 & 2 & 2 \\ 2 & 2 & 2 \\ 2 & 2 & 2 \end{bmatrix}, B = \begin{bmatrix} 0 & 1 & 0 \\ 1 & 0 & 1 \\ 0 & 1 & 0 \end{bmatrix}$$

Apply a **1×1 convolution with 2 filters** (so output depth = 2).

Each 1×1 filter has **3 weights** (one per input channel).

**Filter A (average-ish):** [0.5 0.25 0.25], *bias* = 0

**Filter B (color contrast):** [1 – 1 0.5], *bias* = 0

# 1x1 Convolution

---

## How a $1 \times 1$ works at one pixel

Take the **center pixel** (row 2, col 2):

$$(R, G, B) = (5, 2, 0)$$

**Filter A output:**  $0.5 \times 5 + 0.25 \times 2 + 0.25 \times 0 = 2.5 + 0.5 + 0 = 3.0$

**Filter B output:**  $1 \times 5 - 1 \times 2 + 0.5 \times 0 = 3.0$

So, at that pixel the new 2-channel vector is **[3.0, 3.0]**.

Because the kernel is  $1 \times 1$  (stride 1, no padding), we do this **independently at every spatial location**—it mixes channels but does **not** look at neighbors, and the spatial size stays  $3 \times 3$ .

# Compute the full output maps

Since  $1 \times 1$  is just a per-pixel linear combination across channels, we can write:

$$Y_A = 0.5R + 0.25G + 0.25B, \quad Y_B = 1 \cdot R - 1 \cdot G + 0.5B$$

$Y_A$  (**from Filter A**)

$$0.5R = \begin{bmatrix} 0.5 & 1 & 1.5 \\ 2 & 2.5 & 3 \\ 3.5 & 4 & 4.5 \end{bmatrix}, \quad 0.25G = \begin{bmatrix} 0.5 & 0.5 & 0.5 \\ 0.5 & 0.5 & 0.5 \\ 0.5 & 0.5 & 0.5 \end{bmatrix}, \quad 0.25B = \begin{bmatrix} 0 & 0.25 & 0 \\ 0.25 & 0 & 0.25 \\ 0 & 0.25 & 0 \end{bmatrix}$$

Add them elementwise:

$$Y_A = \boxed{\begin{bmatrix} 1.00 & 1.75 & 2.00 \\ 2.75 & 3.00 & 3.75 \\ 4.00 & 4.75 & 5.00 \end{bmatrix}}$$

$Y_B$  (**from Filter B**)

$$R - G = \begin{bmatrix} -1 & 0 & 1 \\ 2 & 3 & 4 \\ 5 & 6 & 7 \end{bmatrix}, \quad 0.5B = \begin{bmatrix} 0 & 0.5 & 0 \\ 0.5 & 0 & 0.5 \\ 0 & 0.5 & 0 \end{bmatrix}$$

Add them:

$$Y_B = \boxed{\begin{bmatrix} -1.0 & 0.5 & 1.0 \\ 2.5 & 3.0 & 4.5 \\ 5.0 & 6.5 & 7.0 \end{bmatrix}}$$

# Final result (shape change)

---

**Input:**  $3 \times 3 \times 3$

**Output:**  $3 \times 3 \times 2$  (the two maps  $Y_A$  and  $Y_B$  stacked)

So, a **1×1 conv with 2 filters** *shrinks channels* from **3 → 2** while keeping the spatial size **3×3 unchanged**.

---

# RCNN: Region-based Convolutional Neural Network

# Motivation for RCNN


---

- Before R-CNN (2013–2014), object detection relied on:
  - **Traditional methods:** sliding windows, HOG, SIFT, and handcrafted features.
    - These methods were slow and had poor accuracy on complex datasets (e.g., PASCAL VOC).
- R-CNN, proposed by **Ross Girshick et al. (CVPR 2014)**, was the first to **combine region proposals with deep CNN feature extraction**, bringing deep learning to detection.

# RCNN

- The model first extracts *region proposals* (candidate bounding boxes), then applies a **CNN** on each region to classify it and refine its bounding box.


## R-CNN: *Regions with CNN features*



1. Input image



2. Extract region proposals (~2k)



3. Compute CNN features

4. Classify regions

# RCNN

---

## 1. Combining Region Proposals + CNN Features

- Earlier CNNs (like AlexNet) worked only on whole images.
- R-CNN applied CNNs to *regions*, bringing **deep features** into **object localization**.

## 2. Decoupling Detection Stages

- R-CNN separated detection into 3 independent parts:
  1. Proposal generation
  2. CNN-based feature extraction
  3. Classification + regression
- This modularity made it flexible but also computationally heavy.

# RCNN

---

## 3. Fine-Tuning

- After pretraining the CNN on ImageNet classification, R-CNN fine-tuned it for detection using proposed regions labeled as positive/negative examples.
- This was one of the first demonstrations of **transfer learning** in object detection.

## 4. Multi-Stage Training

- R-CNN required **three separate training stages**:
  1. Train CNN on region proposals.
  2. Train SVMs for classification.
  3. Train bounding box regressors.
- These were done independently — time-consuming and complex.

# R-CNN Architecture

---

## **Step 1 — Input Image**

Start with the full image (e.g., 600×1000 pixels).

# R-CNN Architecture

---

## Step 2 — Region Proposal Generation

- Use a **Selective Search** algorithm (not deep learning!) to generate around **2000 candidate regions (RoIs)** likely to contain objects.
  - It merges pixels based on color, texture, and edges.
  - Outputs bounding boxes of varying sizes/aspect ratios.
- *These are "object proposals."*

# R-CNN Architecture

---

## Step 3 — Region Warping

- Each region proposal is:
  - Cropped from the original image.
  - **Warped (resized)** to a fixed size (e.g.,  $227 \times 227$  for AlexNet).
- This ensures each region can be fed into the CNN.

# R-CNN Architecture

---

## Step 4 — Feature Extraction (CNN)

- Each warped region is passed through a **pre-trained CNN** (like AlexNet, VGG16, etc.) to extract a **feature vector**.
  - The CNN is used only **as a feature extractor**.
  - Usually, the output is taken from one of the **fully connected layers (e.g., FC7)** → a **4096-dimensional vector**.
- So, for 2000 regions → 2000 feature vectors.

# R-CNN Architecture

---

## **Step 5 — Classification (SVM)**

- The CNN features for each region are then classified using **class-specific SVMs**.
  - One SVM per object class (e.g., dog, cat, car, etc.).
  - SVM decides whether the region contains that object or background.
- This step replaces the softmax classifier of the CNN.

# R-CNN Architecture

---

## Step 6 — Bounding Box Regression

- For each detected region, a **bounding box regressor** fine-tunes the coordinates of the box to better match the object boundaries.
- This regression is trained separately using the CNN features.

# R-CNN Architecture

---

## **Step 7 — Non-Maximum Suppression (NMS)**

- Multiple boxes may predict the same object → NMS keeps the highest-confidence one and suppresses overlapping duplicates.

# Architecture Summary

---

Stage	Component	Description
1	Selective Search	Generate ~2000 object proposals
2	Warp each region	Resize each to fixed size
3	CNN (AlexNet/VGG)	Extract deep features (e.g., 4096-d)
4	SVM Classifiers	Classify region features into object classes
5	Bounding Box Regressor	Adjust region coordinates
6	NMS	Remove redundant boxes

---

Thank you