



Welcome to Computer Vision



Computer Vision

Dr. Muhammad Tahir

DEPARTMENT OF COMPUTER SCIENCE,
FAST-NUCES, Peshawar

Course Details

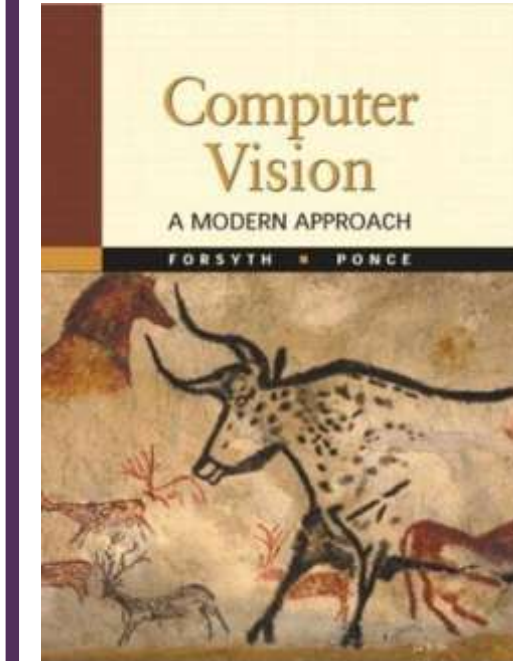
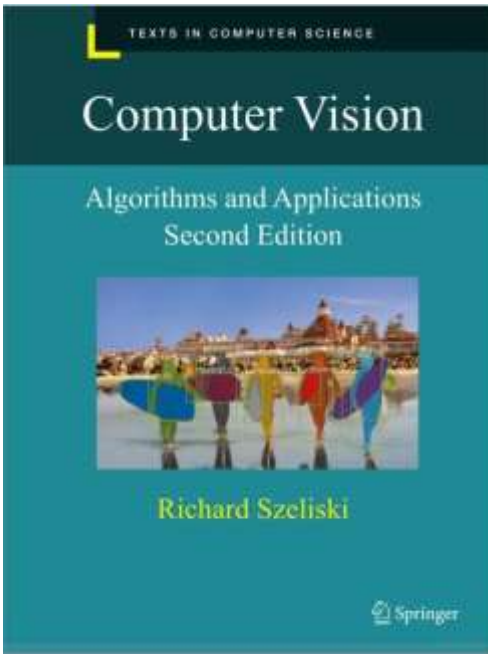
LECTURES: Tuesday
& Wednesday

TIMINGS:
8:00 am – 9:30 am

MY OFFICE:

OFFICE HOURS:

EMAIL: m.tahir@nu.edu.pk



References

The material in these slides are based on:

1

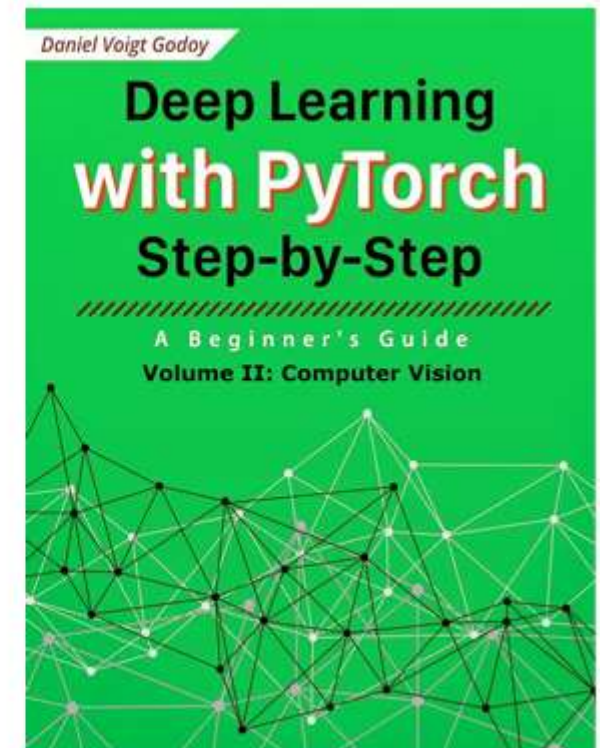
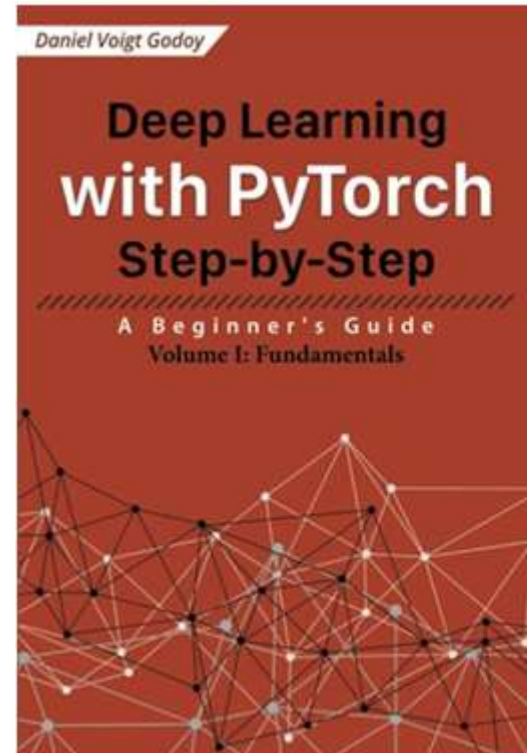
Rick Szeliski's book: [Computer Vision: Algorithms and Applications](#)

2

Forsythe and Ponce: [Computer Vision: A Modern Approach](#)

Recommended Books

Deep Learning with PyTorch Step-by-Step by Daniel Voigt Godoy



Course Learning Outcomes

No	CLO (Tentative)	Domain	Taxonomy Level	PLO
1	Understanding basics of Computer Vision: algorithms, tools, and techniques	Cognitive	2	
2	Develop solutions for image/video understanding and recognition	Cognitive	3	
3	Design solutions to solve practical Computer Vision problems	Cognitive	3	



Outline

Image Representation

Image Filtering

Gradients

Convolution

Correlation

Image Representation

- A digital image is composed of M *rows and N columns* of pixels each storing a value
- Pixel values are most often grey levels in the *range 0-255 (black-white)*
- We will see later that images can easily be represented as matrices

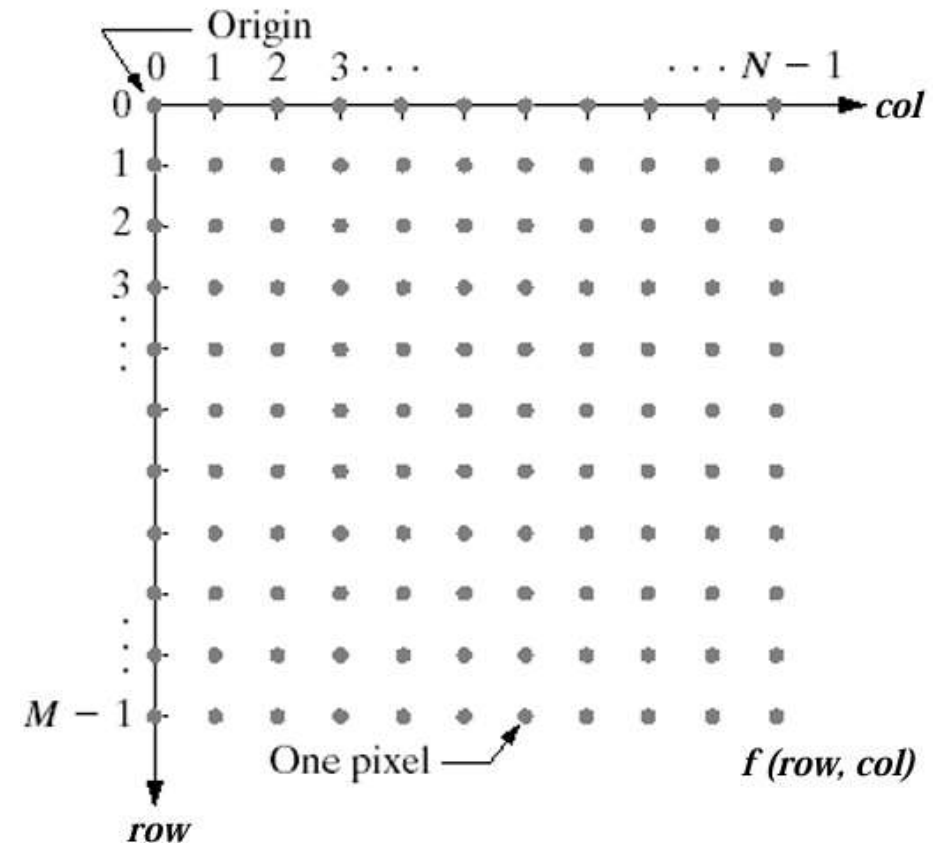
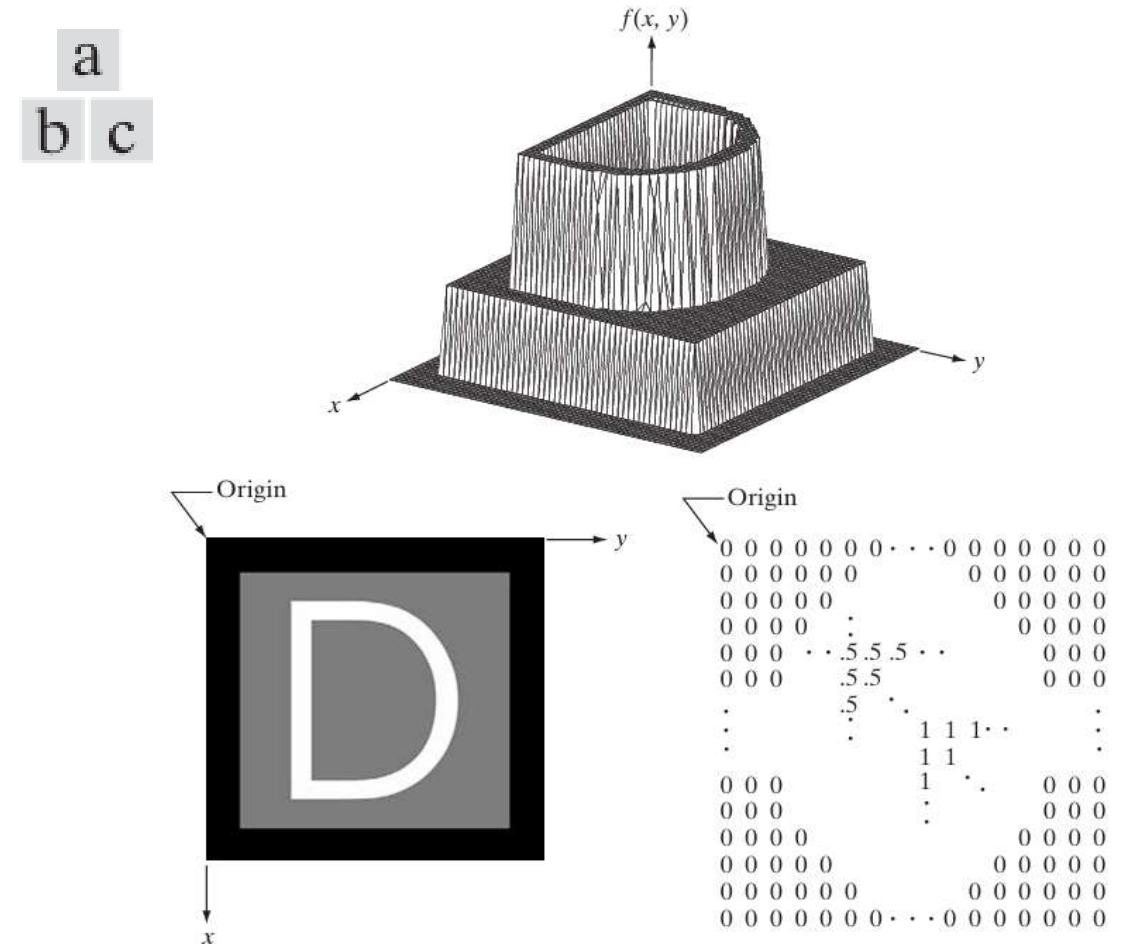


Image Representation

FIGURE 2.18 (a) Image plotted as a surface. (b) Image displayed as a visual intensity array. (c) Image shown as a 2-D numerical array (0,.5,and 1 represent black, gray, and white, respectively).



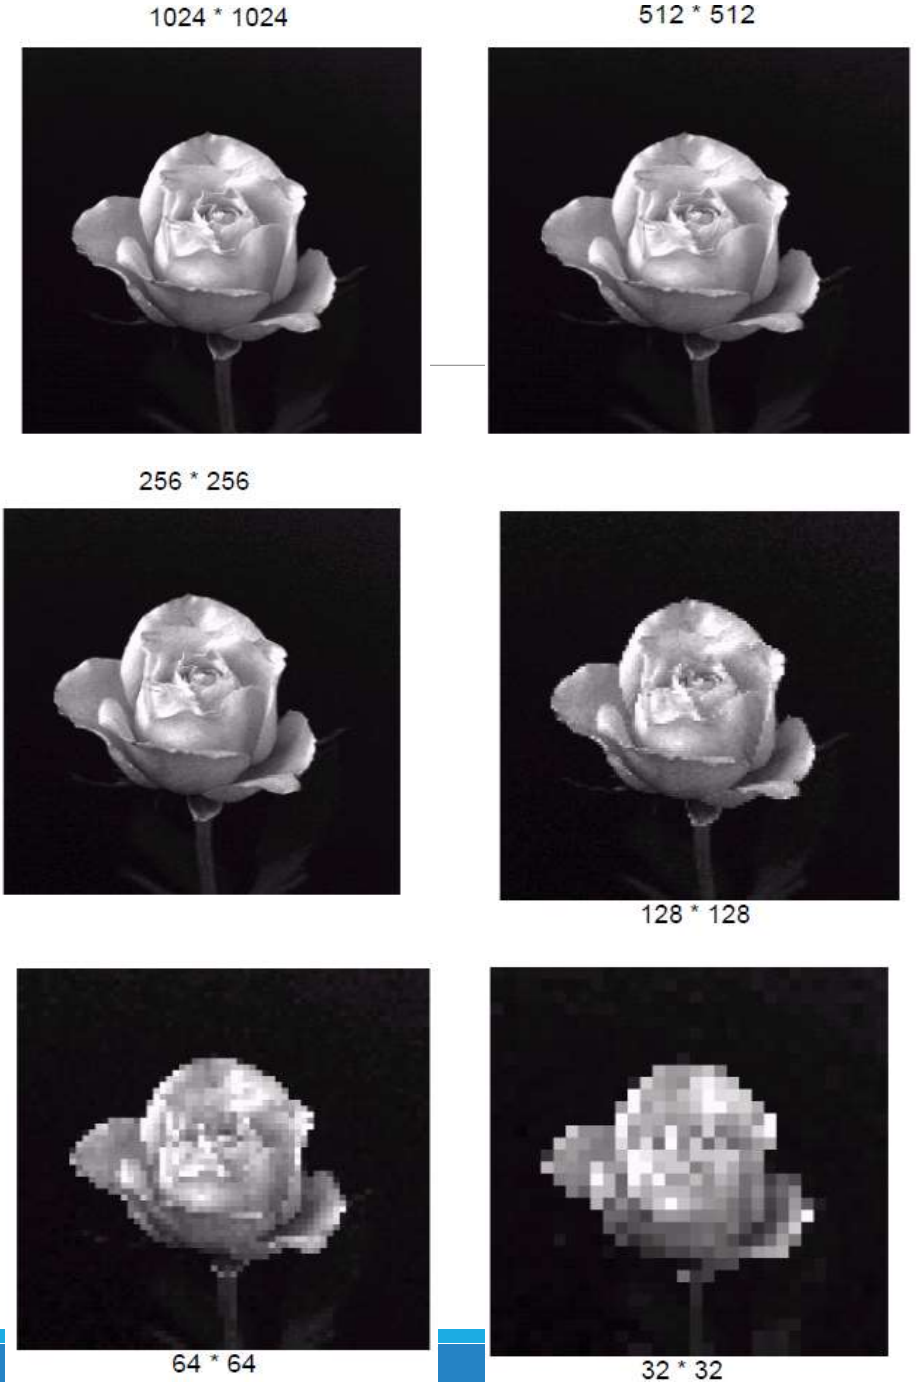
Spatial Resolution

The spatial resolution of an image is determined by how sampling was carried out

Spatial resolution simply refers to the smallest discernible detail in an image

Vision specialists will often talk about pixel size

Graphic designers will talk about dots per inch (DPI)



Intensity Level Resolution

- Intensity level resolution refers to the number of intensity levels used to represent the image
- The more intensity levels used, finer the level of detail discernible in an image
- Intensity level resolution is usually given in terms of the number of bits used to store each intensity level

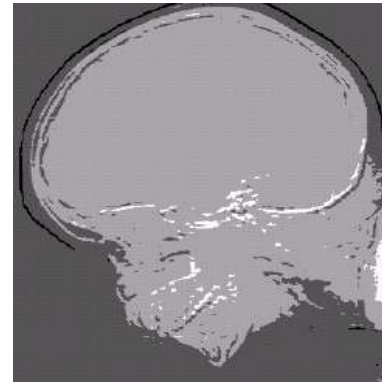
Number of Bits	Number of Intensity Levels
1	2
2	4
4	16
8	256
16	65,536



256 grey levels
(8 bits per pixel)



16 grey levels
(4 bits per pixel)



4 grey levels
(2 bits per pixel)



1 grey levels
(1 bit per pixel)

Resolution: How much is enough?

The big question with resolution is always how much is enough?

This all depends on what is in the image and what you would like to do with it

Key questions include

Does the image look aesthetically pleasing?

Can you see what you need to see within the image?



The picture on the right is fine for counting the number of cars, but not for reading the number plate

Intensity Level Resolution

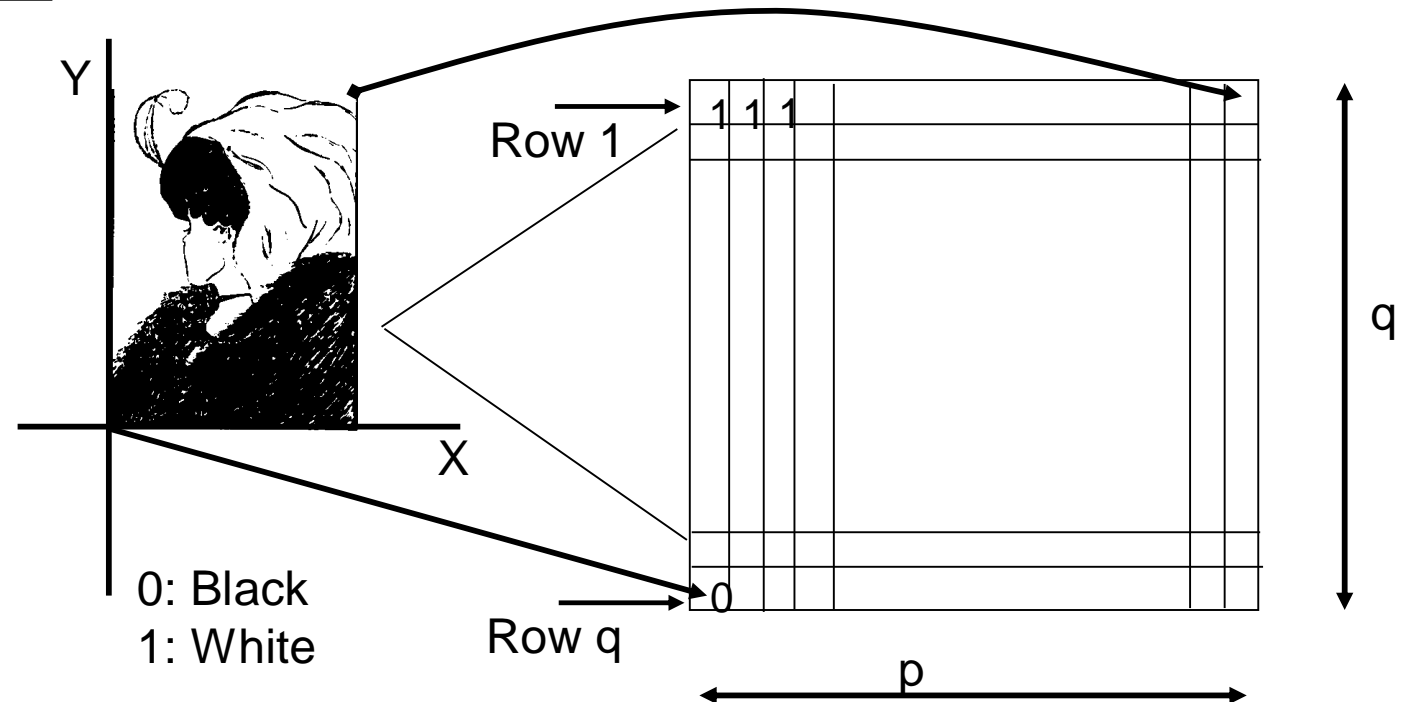
•Based on the **intensity levels** and **number of channels**, we can

categorize images into:

- ✓ Binary image
- ✓ Gray-scale (or gray-tone) image
- ✓ Vector-Valued Images

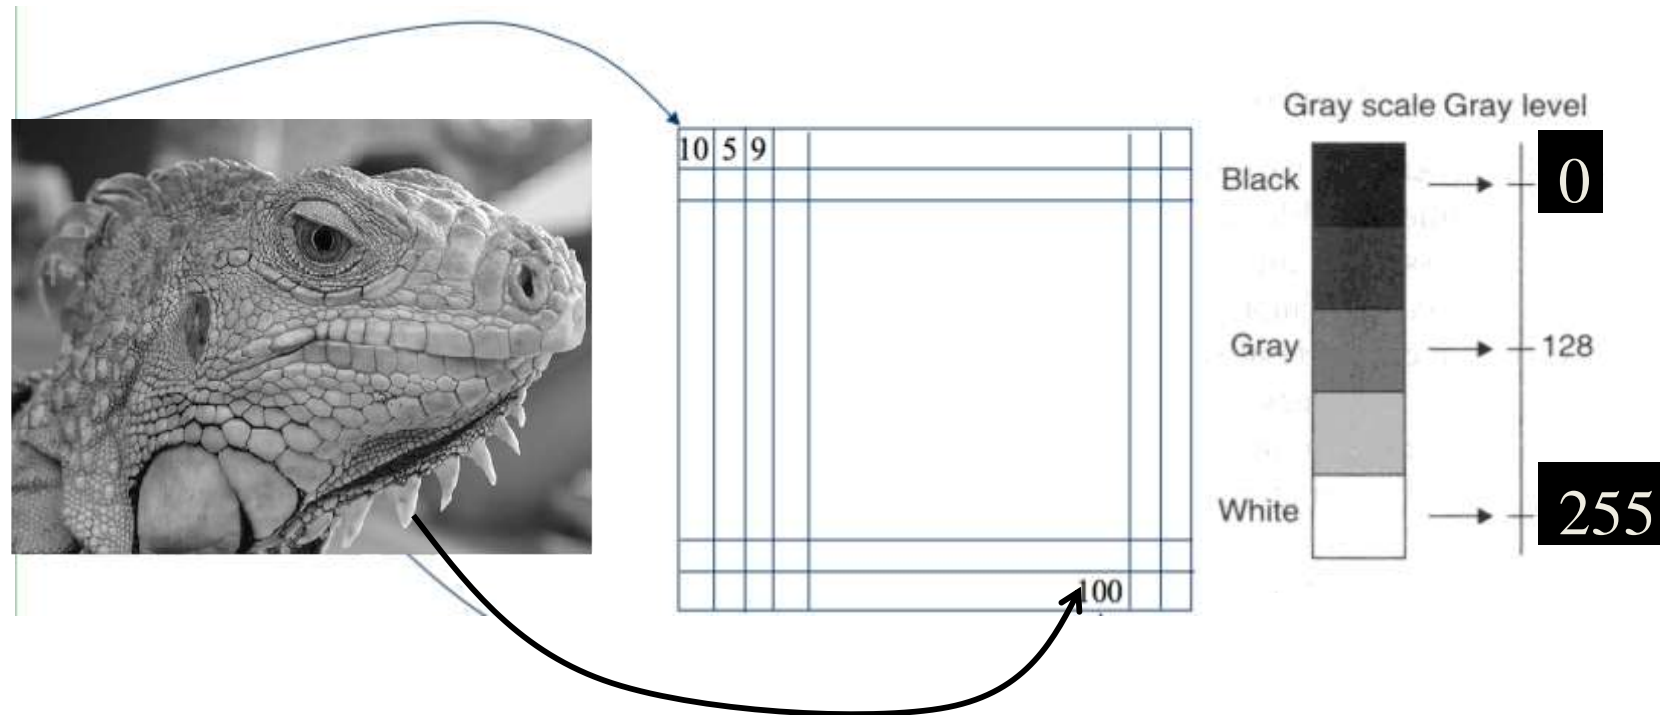
Binary Image

A **binary image** is a digital image that has only two values (**0 or 1**) for each pixel, **0 = black** and **1 = white**



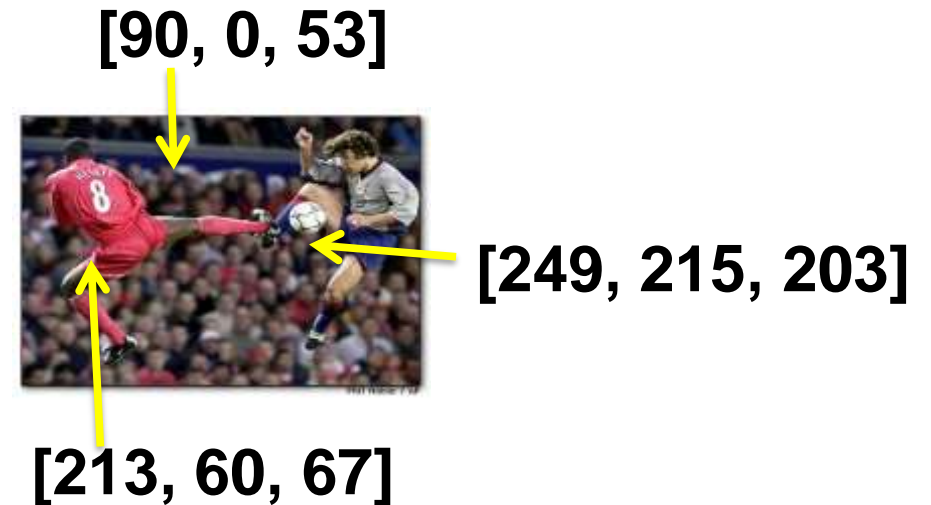
Gray Scale Image

The intensity of each pixel is represented by a value between **0 and 255** (8 bits/pixel)



Vector-Valued Image

- A **vector-valued image** has more than one channel or band
- Example: Color images in the RGB color model have channels for the **red**, **green**, and **blue** component.
- Each channel has a value range from 0 to 255 (like a gray-value image)



RGB Color Images - Example



Red



Green



Blue

Image Enhancement

- Process an image so that the **result** will be more **suitable** than the original image for a specific application.
- The suitability is up to each **application**.
- A method which is quite useful for enhancing an image may not necessarily be the best approach for enhancing another image
- Common reasons for enhancement include
 - ✓ Improving visual quality,
 - ✓ Improving recognition accuracy,
 - ✓ Highlighting interesting detail in images
 - ✓ Removing noise from images

Image Enhancement Methods

- There are two broad categories of image enhancement techniques
 - **Spatial domain Techniques:** Direct manipulation of pixel in an image (on the image plane)
 - **Frequency domain Techniques:** Processing the image based on modifying the Frequency transform of an image
- Many techniques are based on various combinations of methods from these two categories
- For the moment we will concentrate on techniques that operate in the **spatial domain**

Intensity Transformations and Spatial Filtering

- **Spatial Domain**
 - Refers to the image plane itself that contains the pixels
- Image processing techniques in this category are based on direct manipulation of pixels in an image
- Two major categories:
 - **Intensity Transformation techniques** (point processing techniques)
 - Operate on single pixels of an image for the purpose of contrast manipulation and image thresholding
 - **Spatial Filtering techniques** (Neighborhood processing techniques)
 - Perform operations like image sharpening, image smoothing etc., by working in a neighborhood of every pixel in an image

Point operation

- Point operation deals with pixel intensity values individually
- Enhancement at any point depends only on the image value at that point.
- The transformed output pixel value does not depend on any of the neighboring pixel value of the input image
- Point Operation Examples:
 - ✓ Image Negative
 - ✓ Contrast Stretching
 - ✓ Thresholding



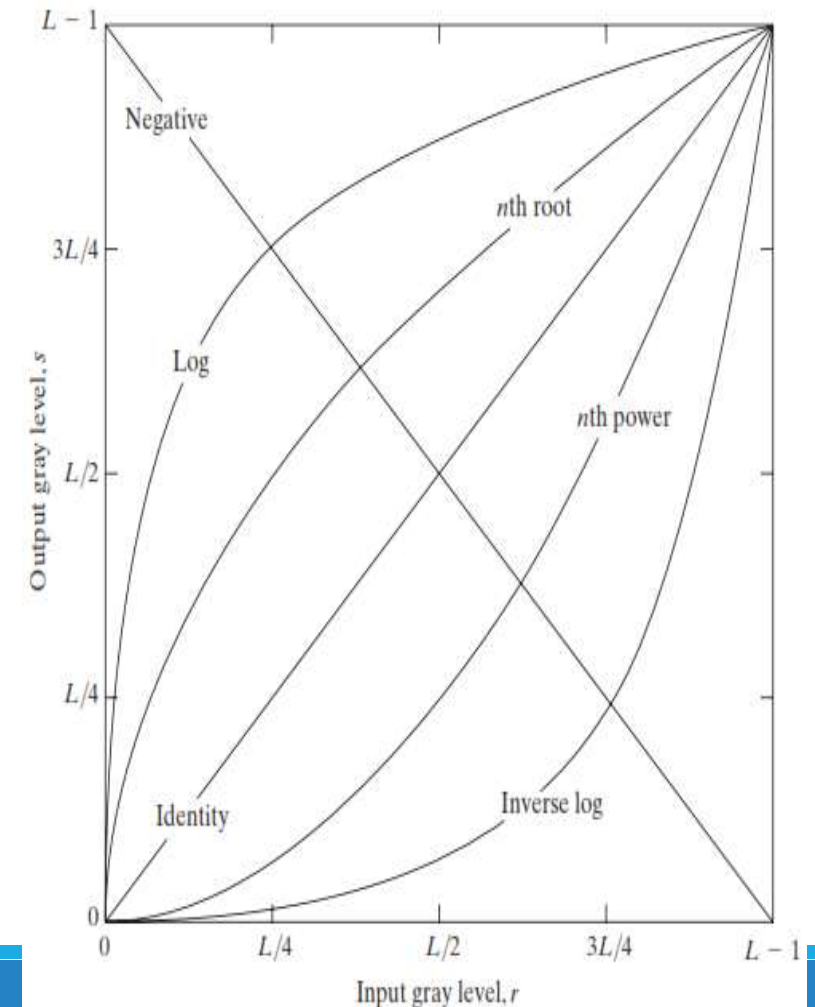
$$g(x,y) = f(x,y) + 20$$



$$g(x,y) = f(-x,y)$$

Point operations: Intensity Transformation

- The intensity values are altered using a particular transformation technique as per the requirement.
- Gray-level transformation function
- **$S=T(r)$**
 - ✓ where r is the pixels of the input image and s is the pixels of the output image.
 - ✓ T is a transformation function that maps each value of r to each value of s
- Basic types of transformation functions that are used frequently in image enhancement:
 - ✓ Linear, Logarithmic: $s = c \log(r + 1)$, Power-Law: $S = C r^\gamma$



Some Basic Intensity Transformation Functions

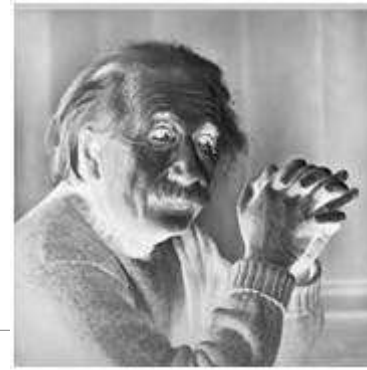
Image Negatives

Log Transformations

Power-Law(Gamma) Transformations

Piecewise-Linear Transformation Functions

Image Negatives



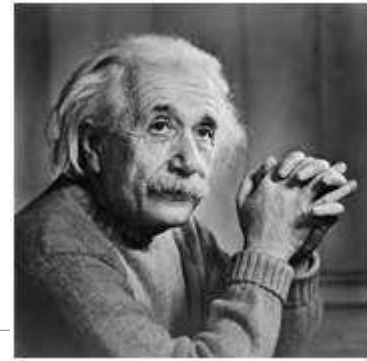
- **Reversing** the intensity levels of an image.

- **Image negative** is produced by **subtracting** each pixel from the **maximum intensity** value. e.g. for an 8-bit image, the max intensity value is $2^8 - 1 = 255$

- An image with gray level in the range $[0, L-1]$ where $L = 2^n$; $n = 1, 2, \dots$

- Negative transformation : $s = L - 1 - r$

Image Negatives



- Reversing the intensity levels of an image.
- Image negative is produced by subtracting each pixel from the maximum intensity value. e.g. for an 8-bit image, the max intensity value is $2^8 - 1 = 255$
- An image with gray level in the range $[0, L-1]$ where $L = 2^n$; $n = 1, 2, \dots$
- Negative transformation : $s = L - 1 - r$

Image Negatives: Applications

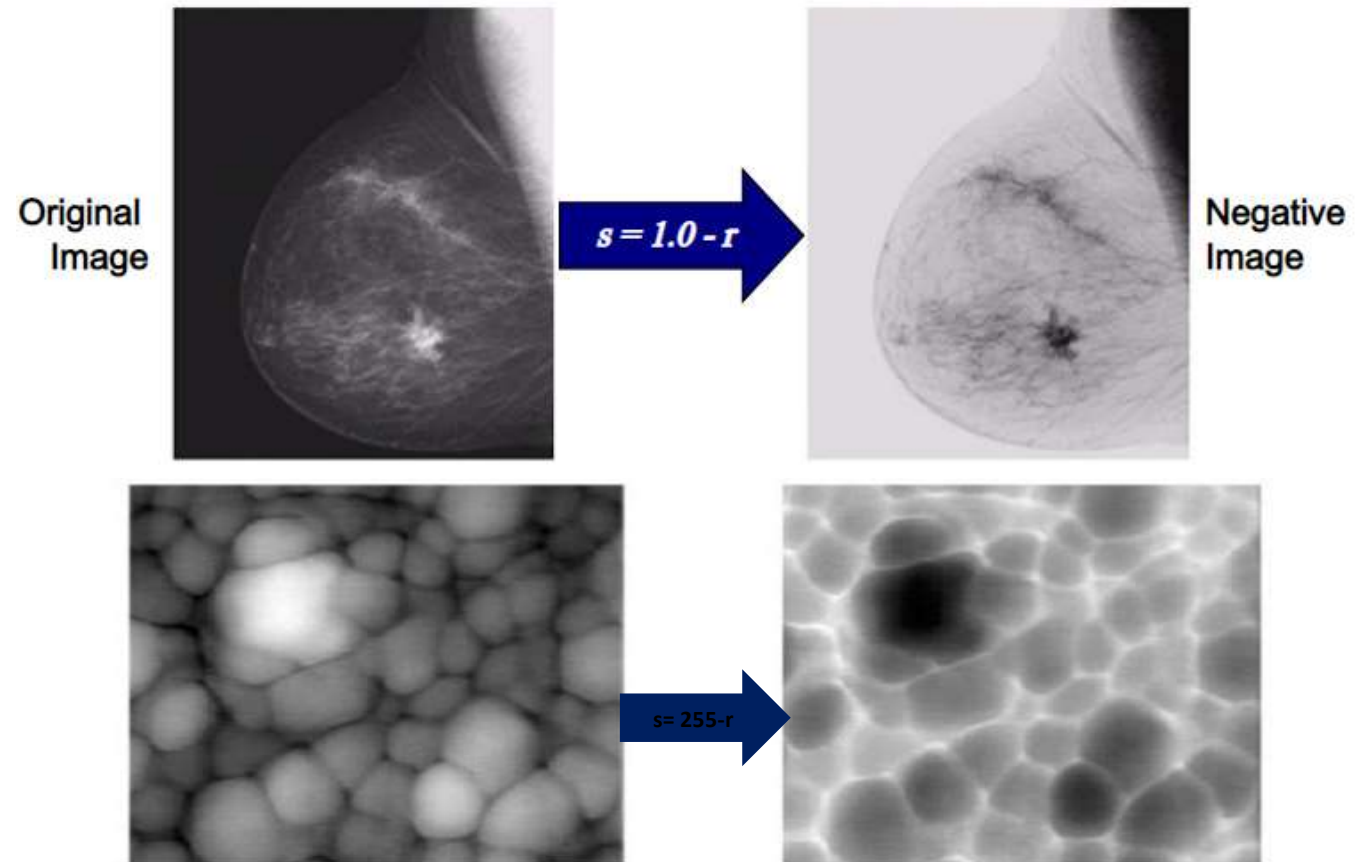


Image Negatives: Activity

An Image with $n = 7$

95	22	36
50	90	6
84	5	13

Image Negative

Image Negatives: Activity

An Image with $n = 7$

95	22	36
50	90	6
84	5	13

Image Negative

32	105	91
77	37	121
43	122	114

Some Basic Intensity Transformation Functions

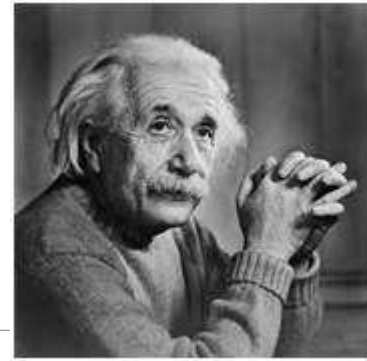
Image Negatives

Log Transformations

Power-Law(Gamma) Transformations

Piecewise-Linear Transformation Functions

Log Transformations

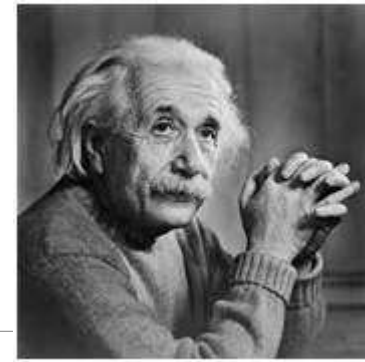


- Log curve maps a **narrow range** of low gray-level values in the input image into a **wider range** of output levels.
- c is a constant and $r \geq 0$
- Used to **expand** the values of **dark pixels** in an image while **compressing** the **higher-level values**.

$$s = c \log(r + 1)$$

Log Transformations

- Log curve maps a **narrow range** of low gray-level values in the input image into a **wider range** of output levels.
- c is a constant and $r \geq 0$
- Used to **expand** the values of **dark pixels** in an image while **compressing** the **higher-level values**.



$$s = c \log(r + 1)$$

Some Basic Intensity Transformation Functions

Image Negatives

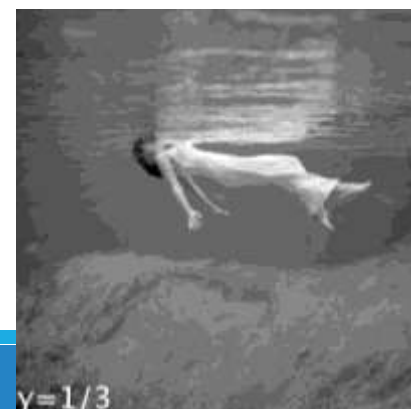
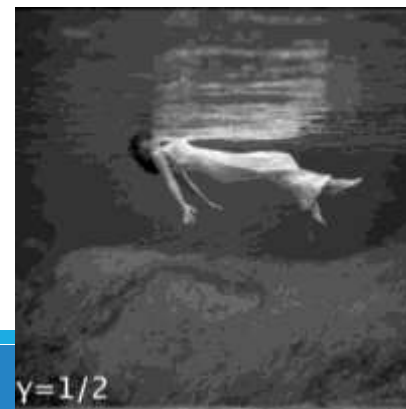
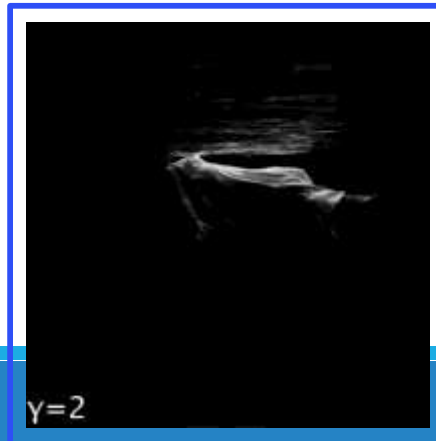
Log Transformations

Power-Law (Gamma) Transformations

Piecewise-Linear Transformation Functions

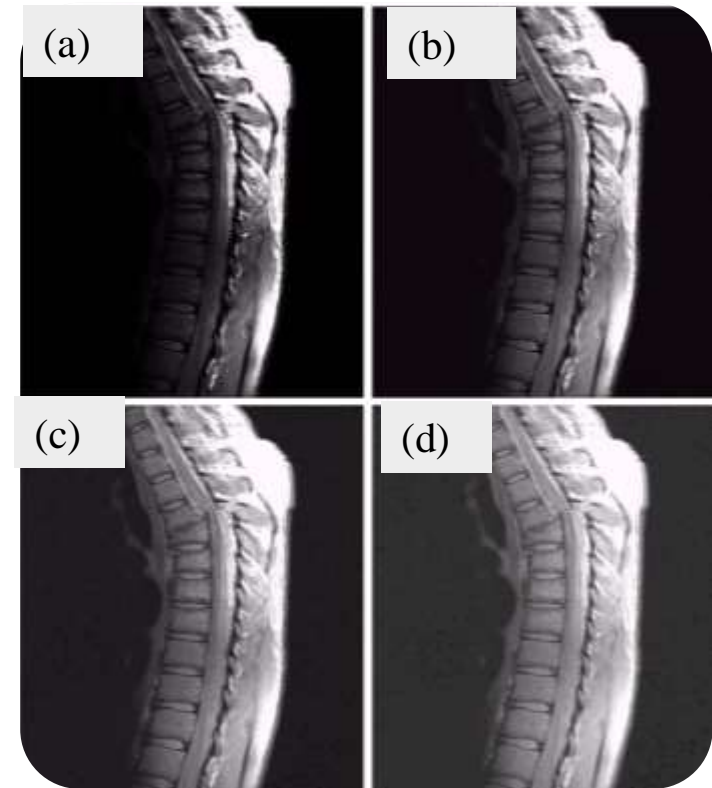
Power-Law (Gamma) Transformations

- Gamma correction function is used to correct image's luminance.
- Power-law transformations: $s = cr^\gamma$ or $s = c(r + \epsilon)^\gamma$
 - ✓ γ : gamma, gamma correction
 - ✓ $\gamma < 1$ maps a **narrow range of dark** input values into a **wider range** of output values [**increase brightness**]
 - ✓ $\gamma > 1$ maps a **narrow range of bright** input values into a **wider range** of output values [**increase darkness**]



Power-Law (Gamma) Transformations

- a) A magnetic resonance image (MRI) of an upper thoracic human spine with a fracture dislocation and spinal cord impingement
- The picture is predominately dark
 - An expansion of gray levels are desirable \Rightarrow needs $\gamma < 1$
- b) result after power-law transformation with $\gamma = 0.6$, $c=1$
- c) transformation with $\gamma = 0.4$ (best result)
- d) transformation with $\gamma = 0.3$ (under acceptable level)



Some Basic Intensity Transformation Functions

Image Negatives

Log Transformations

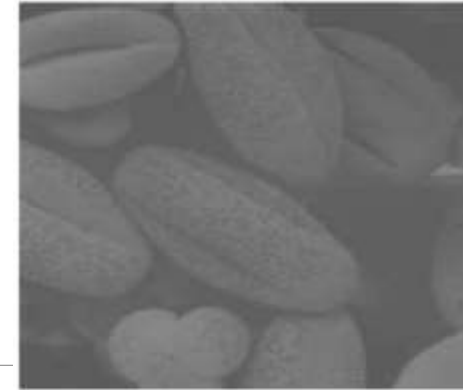
Power-Law (Gamma) Transformations

Piecewise-Linear Transformation Functions

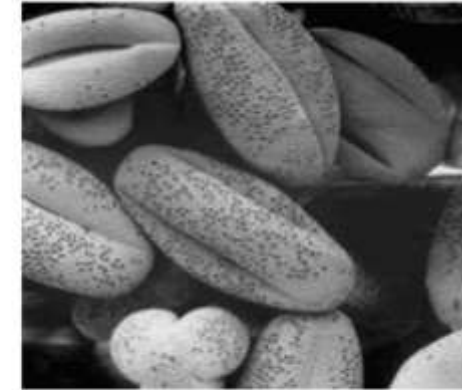
Piecewise-Linear Transformation Functions

- Some commonly used piece-wise linear transformations are:
 - ✓ Contrast Stretching
 - ✓ Gray-level Slicing
 - ✓ Clipping

Contrast Stretching



Original Image

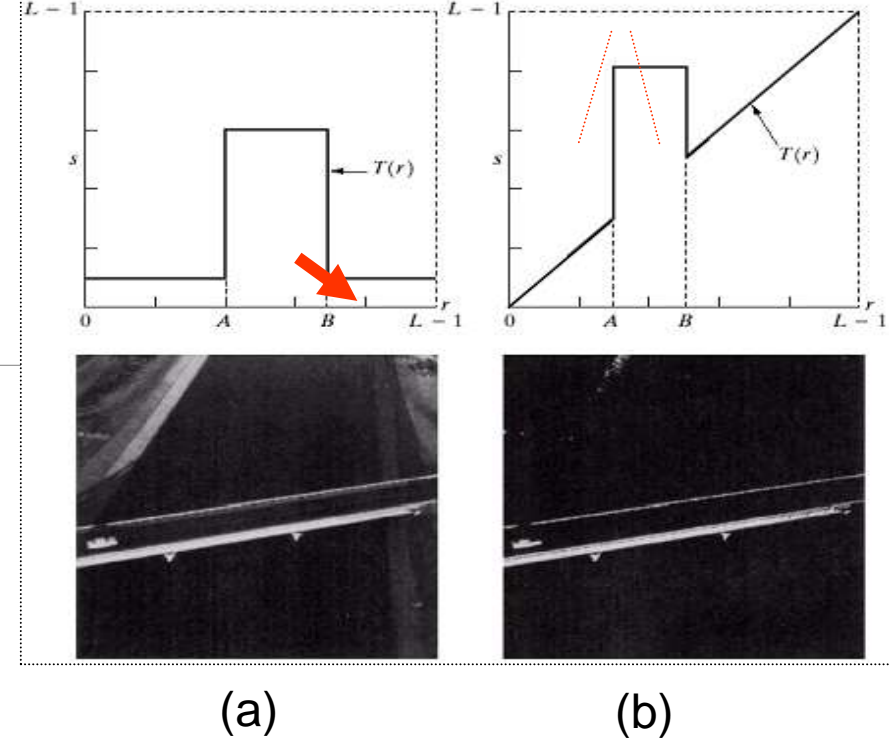


Contrast Enhanced Image

- It is the difference between the intensity values of darker and brighter pixels
- Contrast stretching expands the range of intensity levels in an image.
- Contrast stretching is done in three ways:
 1. Multiplying each input pixel intensity value with a constant scalar.
 - Example: $s=2*r$
 2. Using Histogram Equivalent
 3. Applying a transform which makes dark portion darker by assigning slope of < 1 and bright portion brighter by assigning slope of > 1 .

Gray-level Slicing

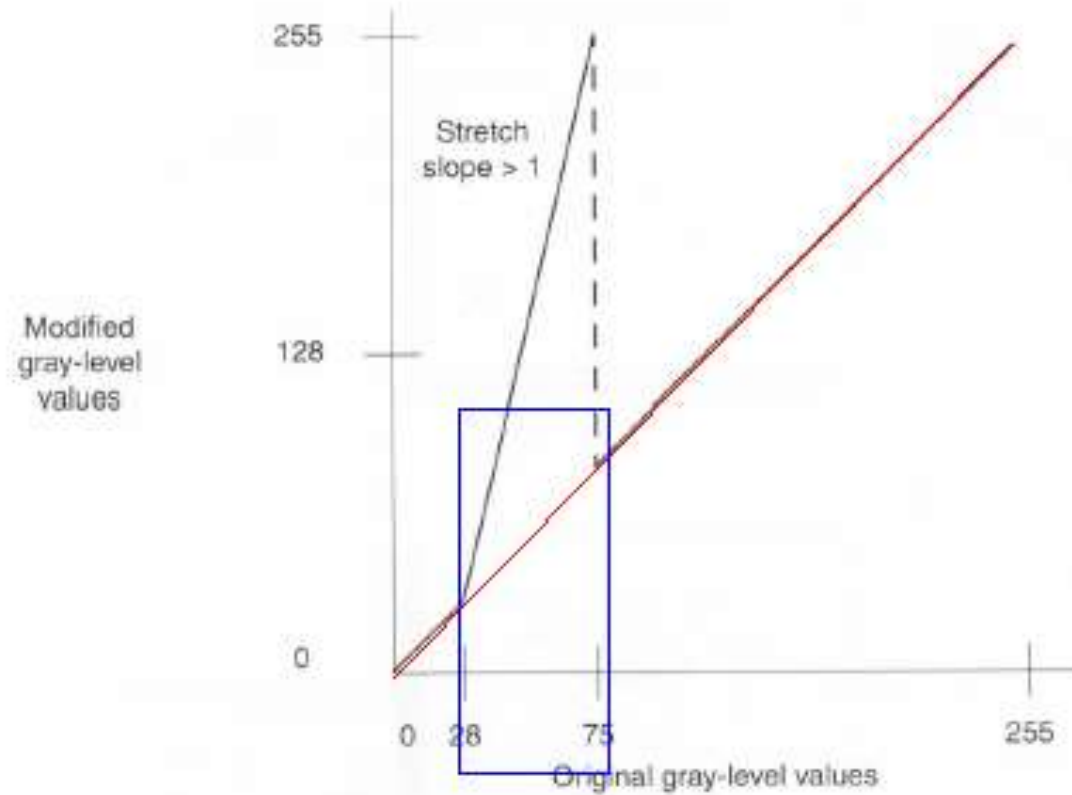
- Highlighting a specific range of gray levels in an image
- Display a high value of all gray levels in the range of interest and a low value for all other gray levels
 - a) transformation highlights range $[A,B]$ of gray level and reduces all others to a constant level
 - b) transformation highlights range $[A,B]$ but preserves all other levels



Gray-level Slicing Example



b. Original image.

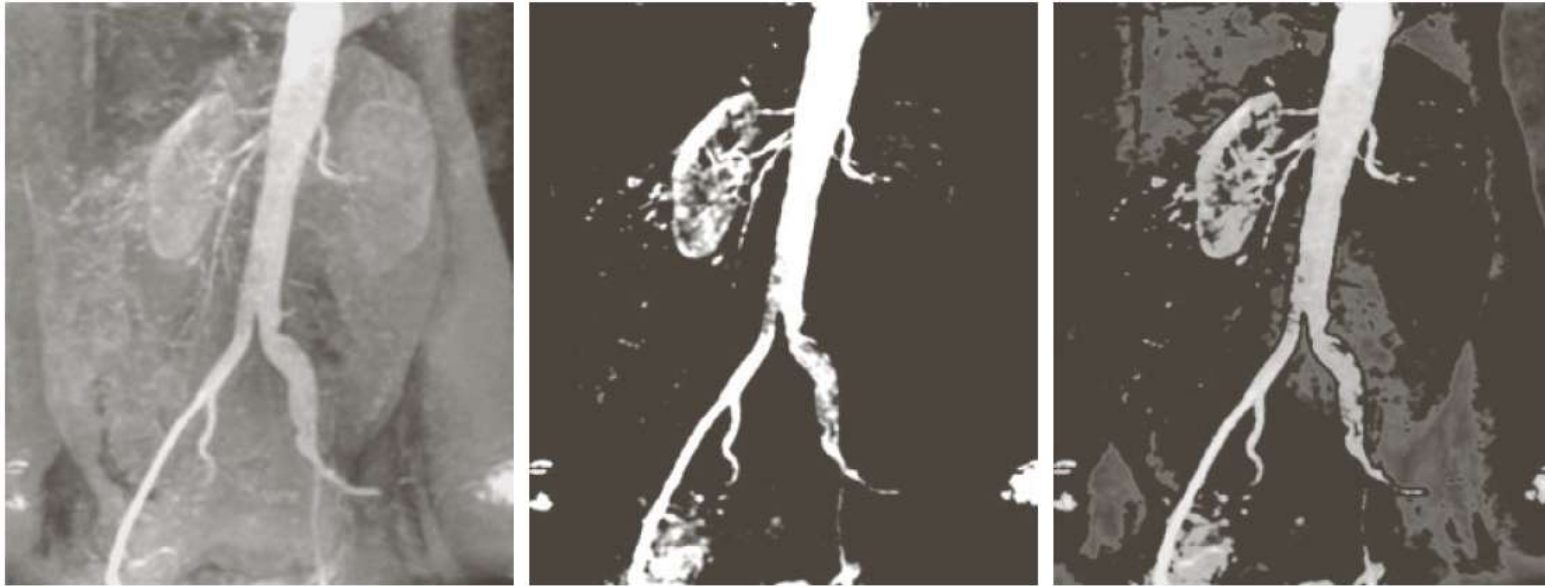


a. Gray-level stretching.



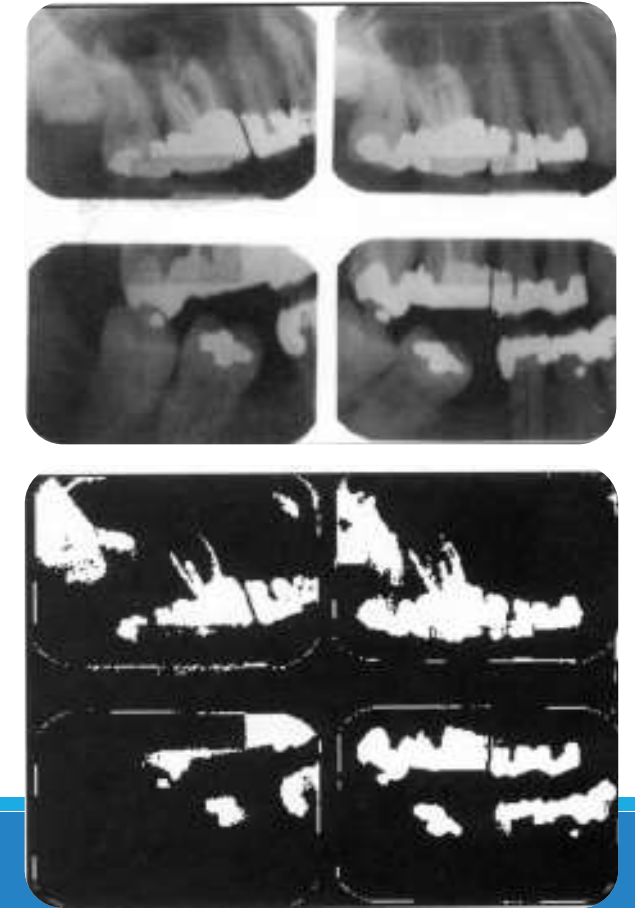
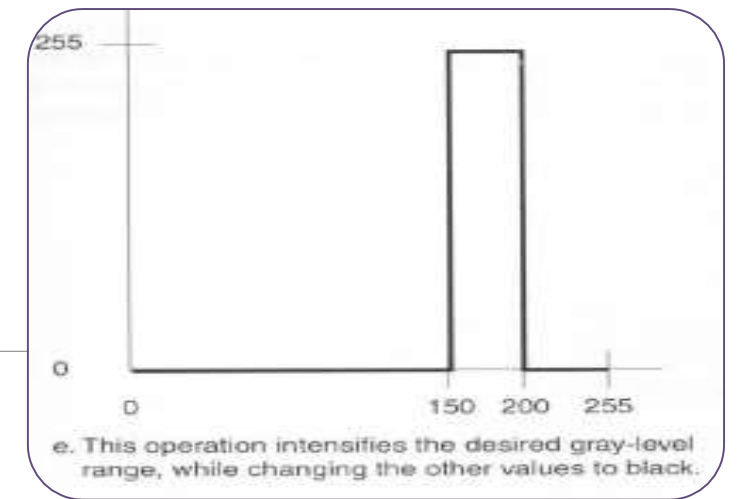
c. Image after modification.

Gray-level Slicing Example



a b c

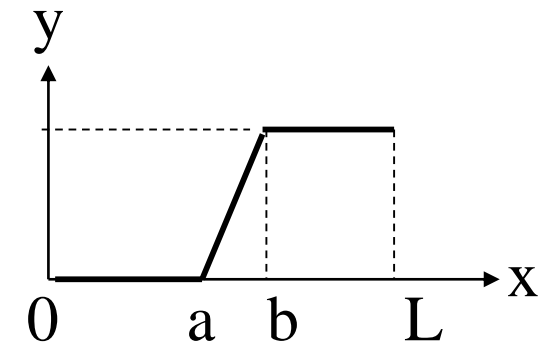
FIGURE 3.12 (a) Aortic angiogram. (b) Result of using a slicing transformation of the type illustrated in Fig. 3.11(a), with the range of intensities of interest selected in the upper end of the gray scale. (c) Result of using the transformation in Fig. 3.11(b), with the selected area set to black, so that grays in the area of the blood vessels and kidneys were preserved. (Original image courtesy of Dr. Thomas R. Gest, University of Michigan Medical School.)



Clipping

- This function truncates all intensities below the optional value Min to Min and all intensities above the optional value Max to Max.
- Clipping can be used to remove **unwanted features**, **noise**, or **extraneous information** from an image.

$$y = \begin{cases} 0 & 0 \leq x < a \\ \beta(x - a) & a \leq x < b \\ \beta(b - a) & b \leq x < L \end{cases}$$



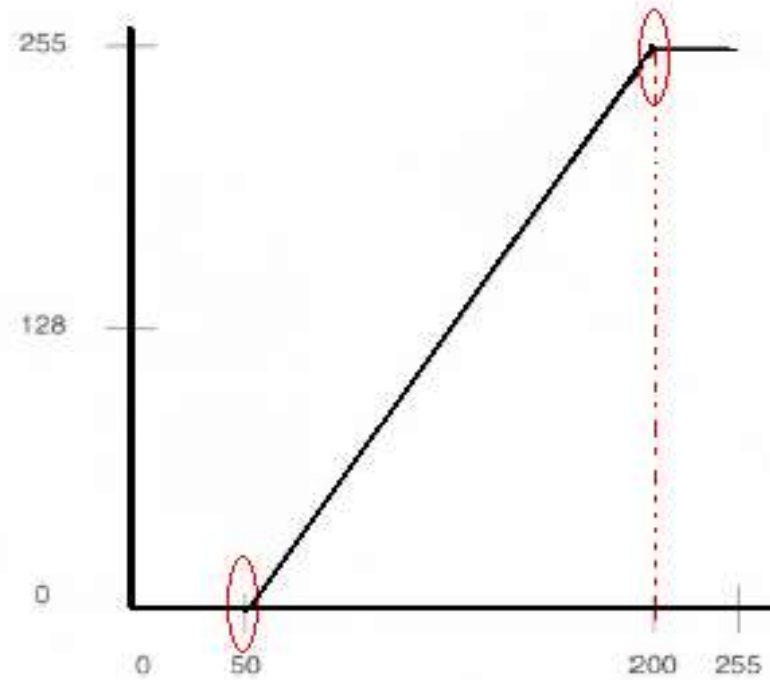
$$a = 50, b = 150, \beta = 2$$

Clipping Example



e. Original image.

Modified gray-level values



f. Image after modification.

Piecewise-Linear Transformation Functions

•**Advantage:**

- ✓ The form of piecewise functions can be arbitrarily complex.
- ✓ A Practical Implementation of some important transformations can be formulated only as piecewise functions.

•**Disadvantage:**

- ✓ Their specification requires considerably more user input.

Filter (Mask) Operations

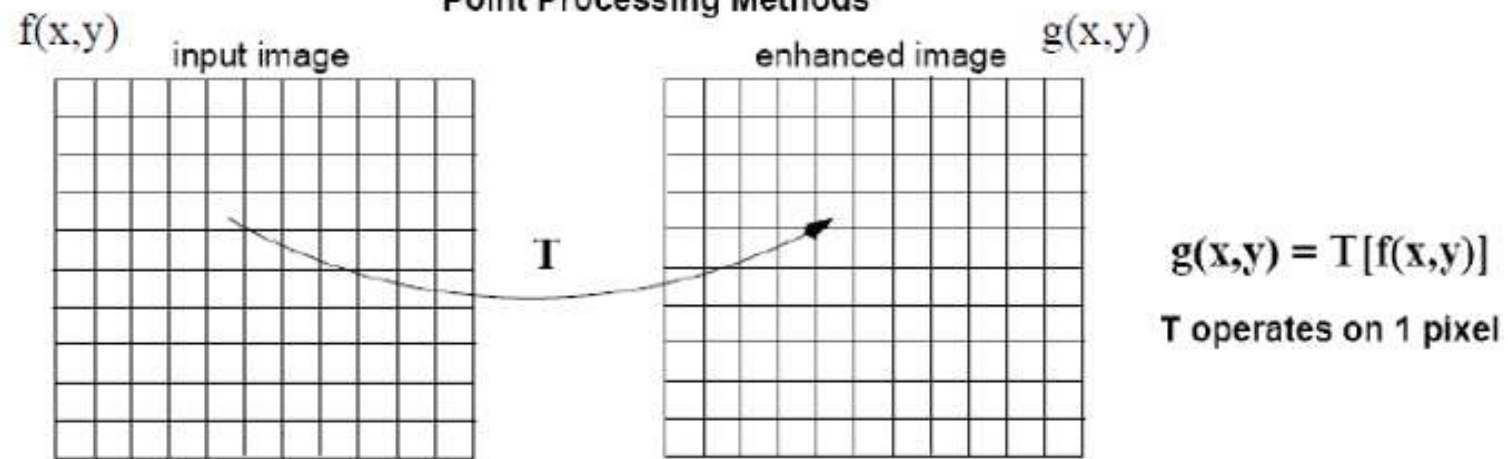
Why Spatial Domain?

- Some image processing tasks are easier, or more meaningful to implement in the spatial domain, while others are best suited for other domains.
- Spatial domain techniques are more efficient computationally and require fewer processing resources to implement.
- Point-operations don't know anything about their neighbors
- Most image features (edges, textures, etc) involve a spatial neighborhood of pixels
- If we want to enhance or manipulate these features, we need to go beyond point operations

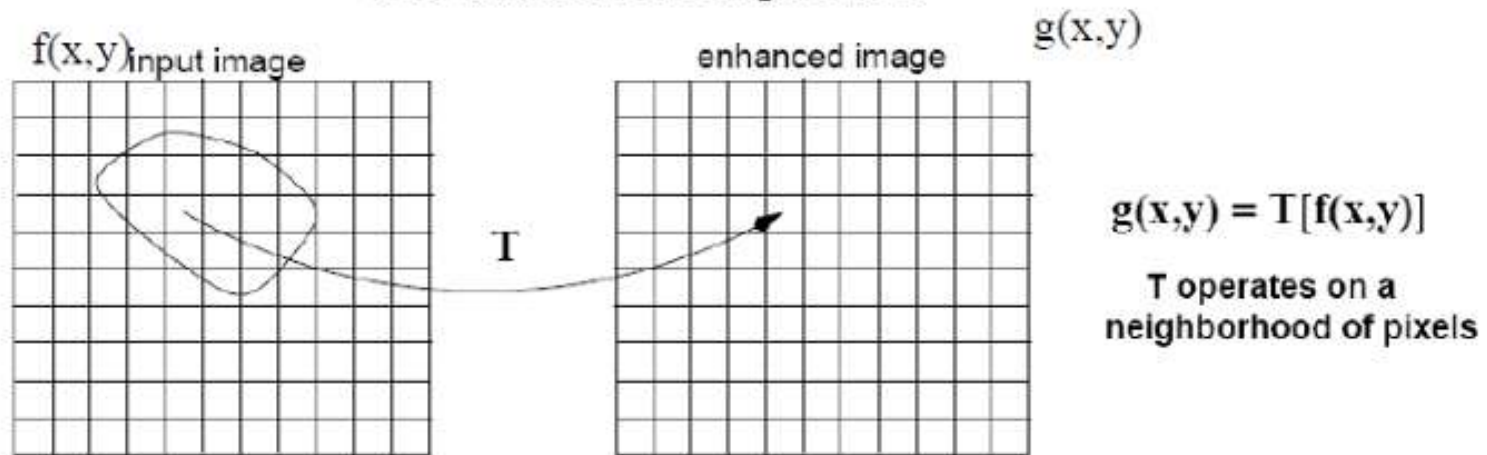
Neighborhood Processing

- To modify or change a pixel value we only need to know its local neighbour(s) gray values.
- Also referred as local processing
- Spatial Filtering is a special case of neighborhood processing

Point Processing Methods

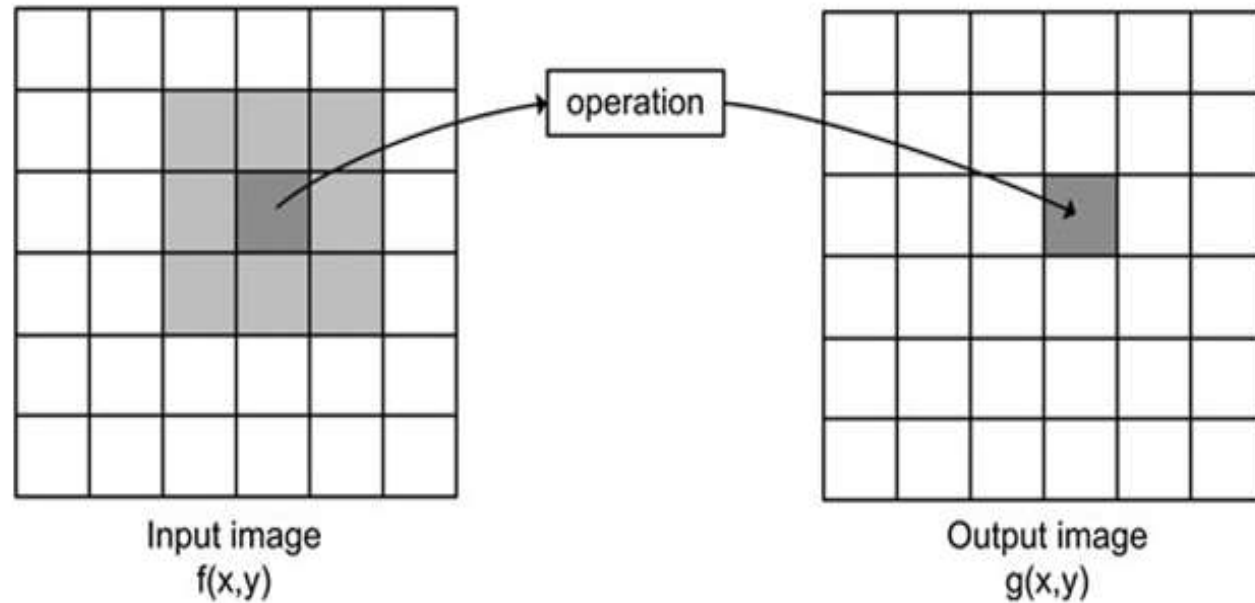


Area or Mask Processing Methods



Spatial Filtering

- Spatial domain filtering
- Capabilities of point operations are limited
- **Filters:** combine pixel's value + values of neighbors



What Point Operations Can't Do

- Combining multiple pixels needed for certain operations:
 - **Enhance an image**, e.g., denoise, Blurring, Sharpening.
 - **Extract information**, e.g., texture, edges.
 - **Detect patterns**, e.g., template matching.



Filtering Operations Use “filters”

36	36	36	36	36
36	36	45	45	45
36	45	45	45	54
36	45	54	54	54
45	45	54	54	54

Input Image

1/9	1/9	1/9
1/9	1/9	1/9
1/9	1/9	1/9

filter

**	**	**	**	**
**	39	**	**	**
**	**	**	**	**
**	**	**	**	**
**	**	**	**	**

Output Image

- **Filters (masks)** operate on a neighborhood of pixels.
- The sub-image is called a **filter** (or mask, kernel, template, window).
- The values in a filter sub-image are referred to as **coefficients**, rather than pixels.
 - ✓ A mask of coefficients is centered on a pixel.
 - ✓ The mask **coefficients** are **multiplied** by the pixel values in its neighborhood, and the products are **summed**.
 - ✓ The result goes into the corresponding pixel position in the output image.

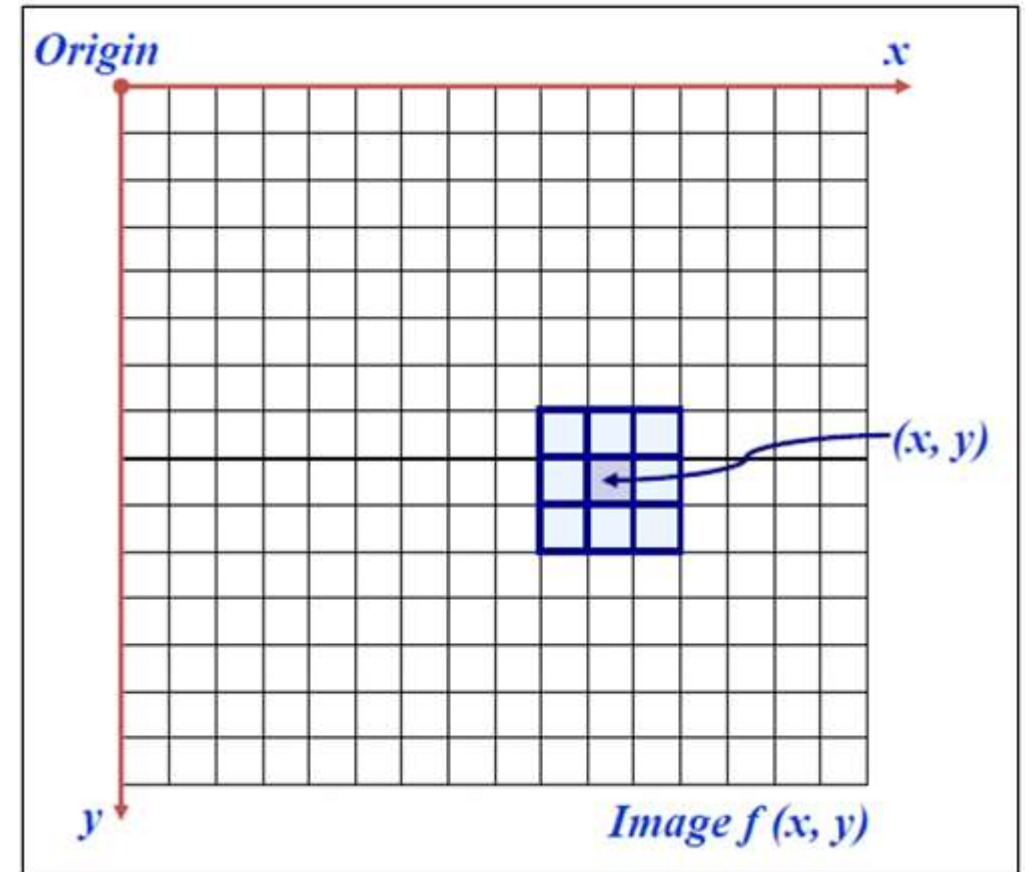
Simple Neighbourhood Operations

- Some simple neighborhood operations include:
 - ✓ **Min:** Set the pixel value to the **minimum** in the neighborhood
 - ✓ **Max:** Set the pixel value to the **maximum** in the neighborhood
 - ✓ **Average:** Set the pixel value to the **average** of the neighborhoods
 - ✓ **Median:** The median value of a set of numbers is the **midpoint** value in that set (e.g., from the set [1, 7, 15, 18, 24] **15** is the median).
Sometimes the median works better than the average

Spatial Filtering

1	-1	-1
-1	1	-1
-1	-1	1

- Move a “mask” (e.g. a square) over the complete image pixel-by-pixel to create a new image
- In this new image, each pixel gray level value is written as the linear combination of the neighbor pixels.
- “Mask” defines the coefficients of linear combination.
 - Mask is also called Filter or Template



Filters

Smoothing Spatial Filters

For blurring and noise reduction

Smoothing Linear Filters

Averaging filter or Low pass filter

Order Statistic (Nonlinear) Filters

Response is based on ordering or ranking of pixels under the kernel

Median Filter

Impulse noise (salt and pepper noise)

Max Filter

Min Filter

FIGURE 3.31
Another
representation of
a general 3×3
filter mask.

w_1	w_2	w_3
w_4	w_5	w_6
w_7	w_8	w_9

a b

FIGURE 3.32 Two 3×3 smoothing (averaging) filter masks. The constant multiplier in front of each mask is equal to 1 divided by the sum of the values of its coefficients, as is required to compute an average.

Averaging Filters

- One of the simplest spatial filtering operations we can perform is a smoothing operation (remove sharp features)
 - Simply average all of the pixels in a neighborhood around a central value
 - Especially useful in removing noise from images
 - Also useful for highlighting gross detail
 - Moving average in 1D:

$$\frac{1}{9} \times \begin{bmatrix} 1 & 1 & 1 \\ 1 & 1 & 1 \\ 1 & 1 & 1 \end{bmatrix}$$

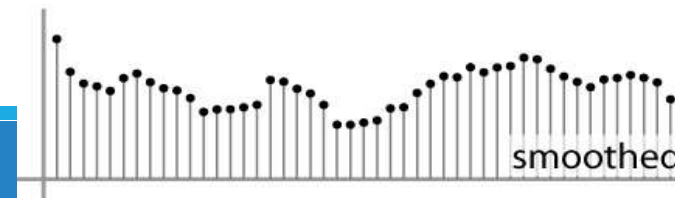
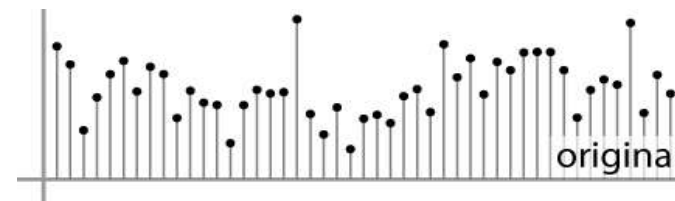
Standard average

$$\frac{1}{16} \times \begin{bmatrix} 1 & 2 & 1 \\ 2 & 4 & 2 \\ 1 & 2 & 1 \end{bmatrix}$$

Weighted average

$$\frac{1}{25} \times$$

1	1	1	1	1
1	1	1	1	1
1	1	1	1	1
1	1	1	1	1
1	1	1	1	1



Average filtering: Example

$$h[\cdot, \cdot] = \frac{1}{9} \begin{bmatrix} 1 & 1 & 1 \\ 1 & 1 & 1 \\ 1 & 1 & 1 \end{bmatrix}$$

 $f[\cdot, \cdot]$ [illegible]
$$g[\cdot, \cdot]$$
[illegible]

Average filtering: Example

$$h[\cdot, \cdot] \quad \frac{1}{9} \begin{bmatrix} 1 & 1 & 1 \\ 1 & 1 & 1 \\ 1 & 1 & 1 \end{bmatrix}$$

 $f[\cdot, \cdot]$ [illegible]
$$g[\cdot, \cdot]$$
[illegible]

Average filtering: Example

$$h[\cdot, \cdot] \quad \frac{1}{9} \begin{bmatrix} 1 & 1 & 1 \\ 1 & 1 & 1 \\ 1 & 1 & 1 \end{bmatrix}$$

$$f[\cdot, \cdot]$$
[illegible]
$$g[\cdot, \cdot]$$
[illegible]

Average filtering: Example

$$h[\cdot, \cdot] = \frac{1}{9} \begin{bmatrix} 1 & 1 & 1 \\ 1 & 1 & 1 \\ 1 & 1 & 1 \end{bmatrix}$$

 $f[\cdot, \cdot]$ [illegible]
$$g[\cdot, \cdot]$$
[illegible]

Average filtering: Example

[illegible][illegible]

Examples: Average Filtering



(a) Original image



(b) Average filtering



(c) Using a 9×9 filter



(d) Using a 25×25 filter



Original image



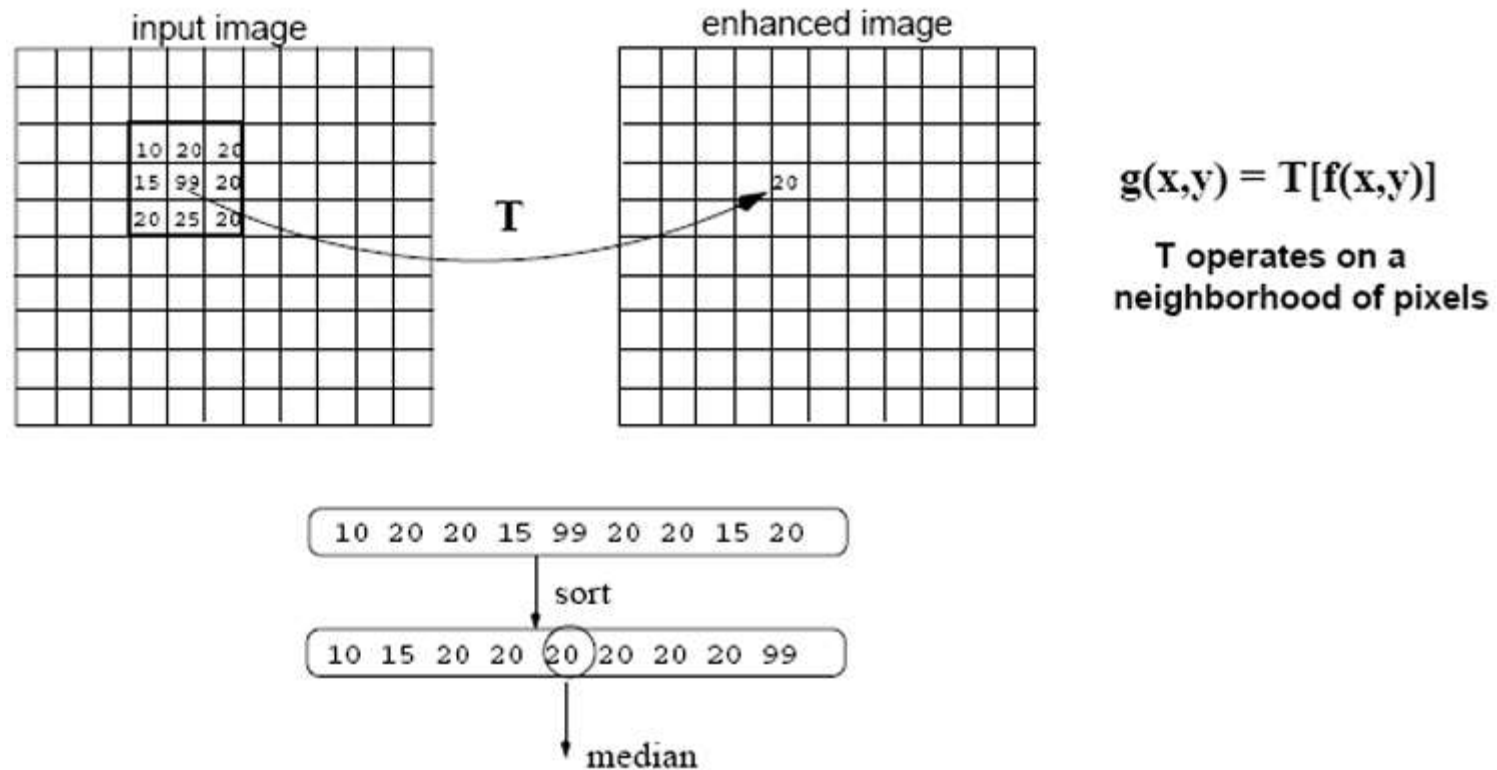
Smoothed image



Very Smoothed image

Median Filtering

Area or Mask Processing Methods



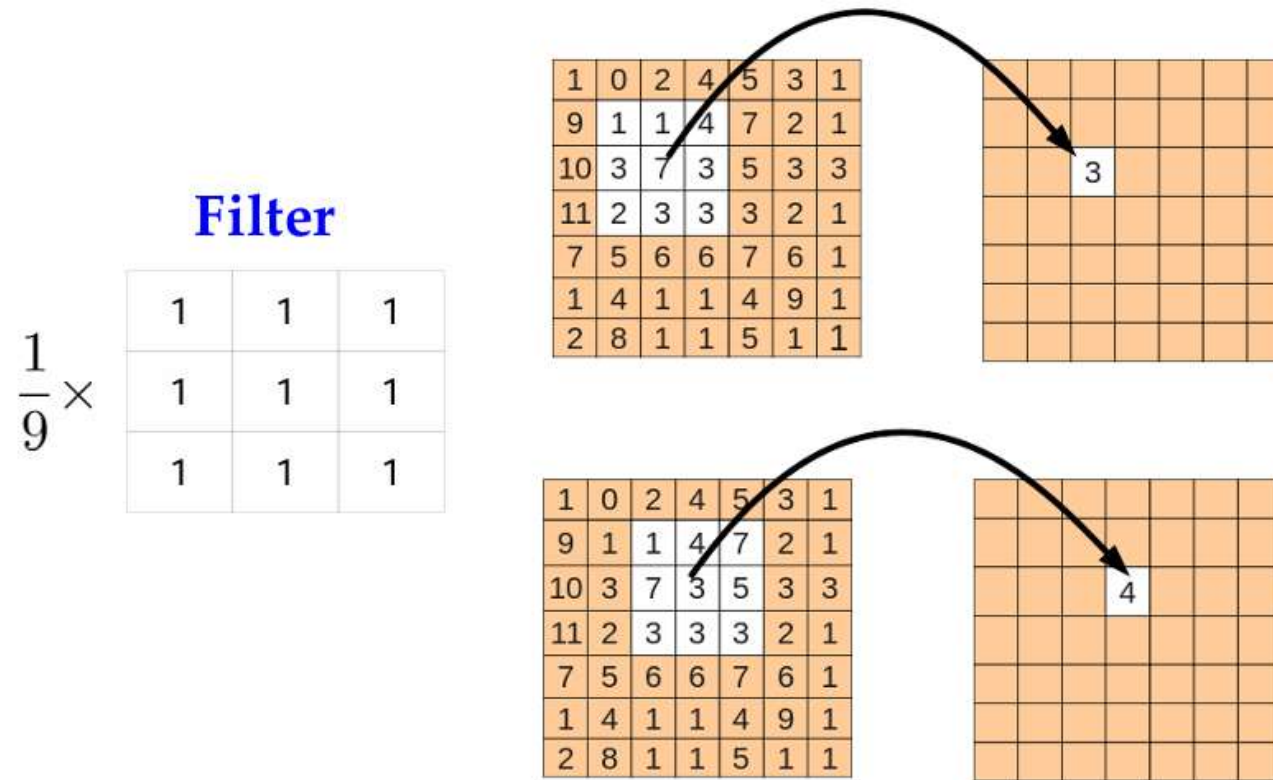
Spatial Filtering Methods

Spatial Filtering Methods

Main linear spatial filtering methods:

- Correlation
- Convolution

Spatial Filtering as Correlation



Thank you

A solid blue horizontal bar spanning the entire width of the slide at the bottom.